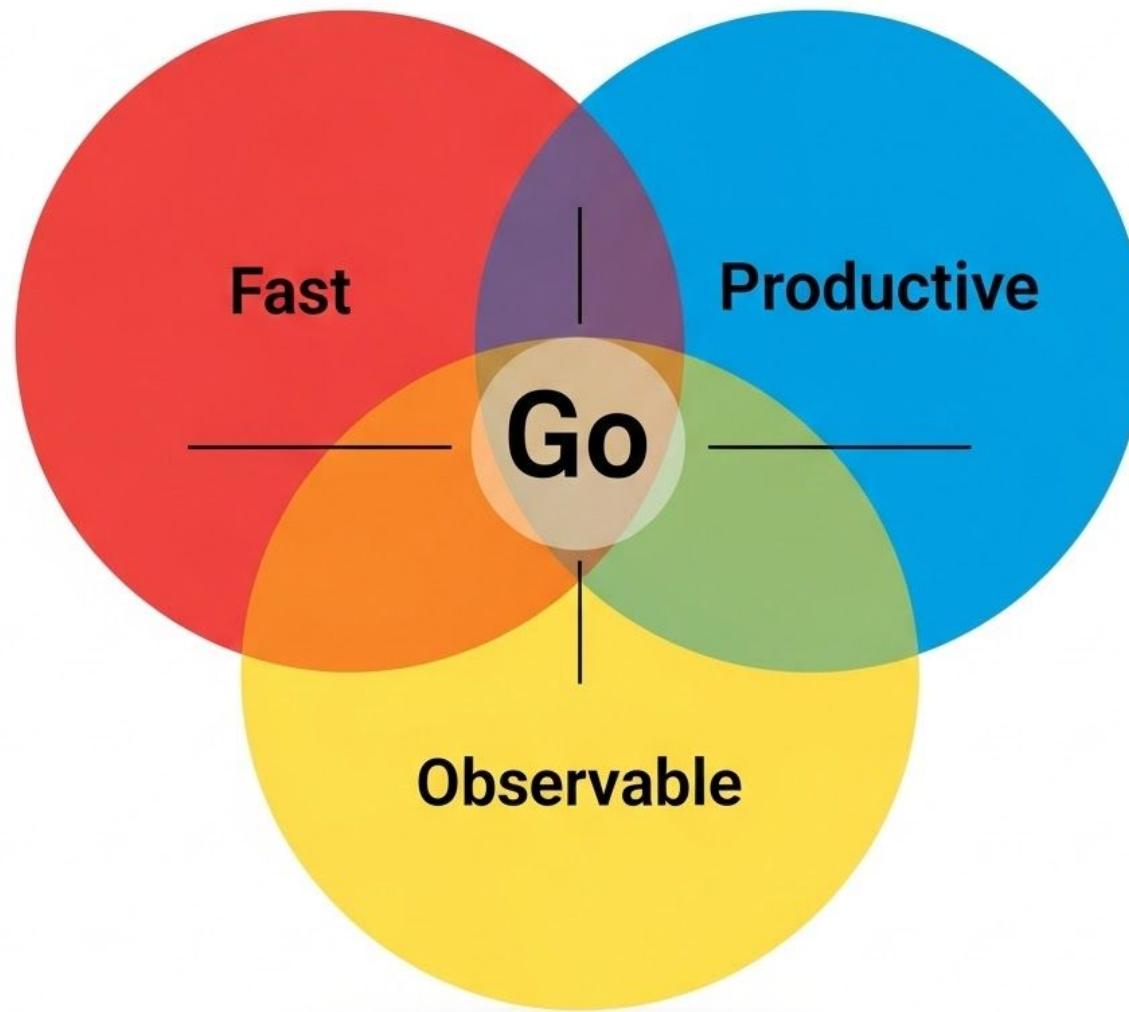


# Profiling Request Latency with Critical Path Analysis

Felix Geisendorfer



# Do you like fast software?



# Go makes Continuous Profiling Easy

- net/http/pprof exposes runtime/pprof and runtime/trace data with < 1-2% overhead

root

errgroup.(\*Group).add.func1

converter.(\*RowConverter).parseRows-fm

converter.(\*RowConverter).parseRows

log.(\*traceLogger).Errorf

log.(\*zapLogger).Error

zap.(\*Logger).Error

zapcore.(\*CheckedEntry).Write

zapcore.(\*ioCore).Write

zapcore.(\*lockedWriteSyncer).Write

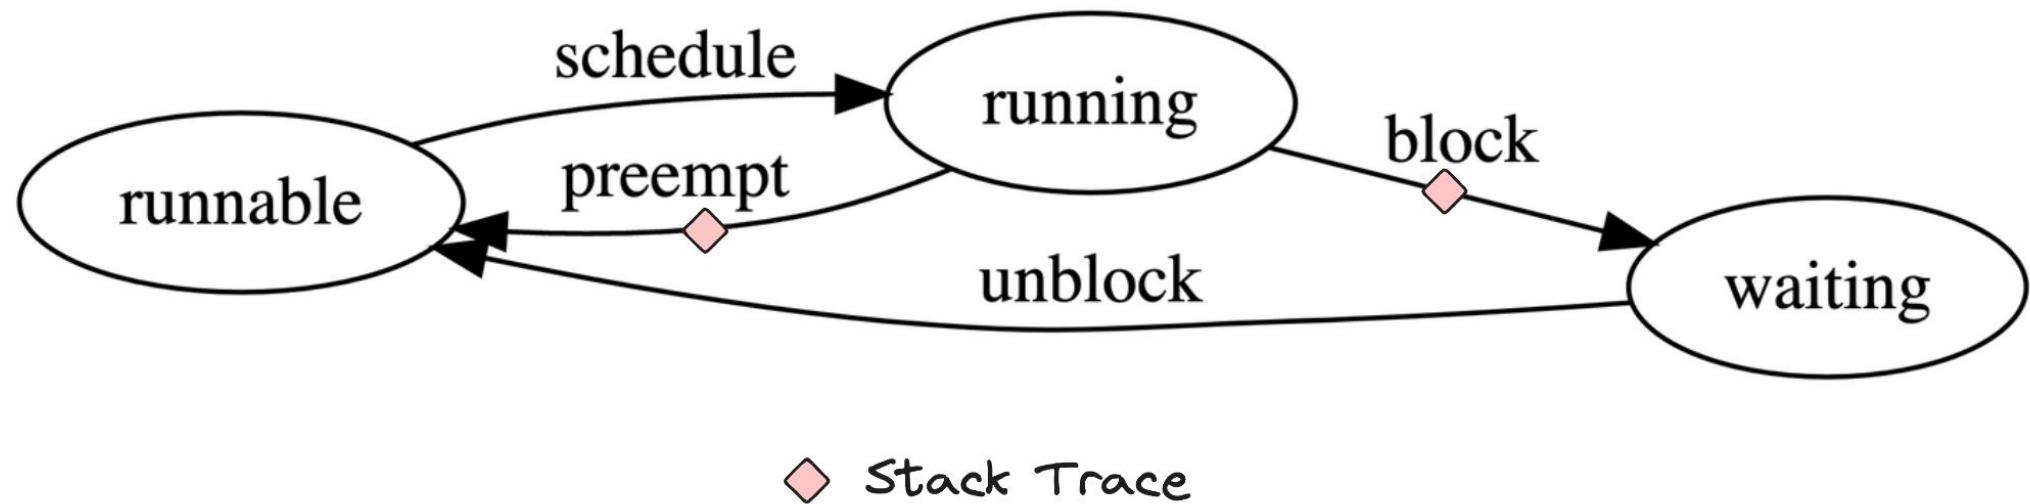
sync.(\*Mutex).Unlock

# For endpoint X, how much is p99 latency caused by:

- CPU Usage
- Mutex Contention
- Channel Contention
- Scheduling Latency
- Garbage Collection (GC Assists)
- Syscalls
- Sleep
- Networking (*prefer distributed tracing for this*)

i.e. a complete breakdown of latency with nothing missing

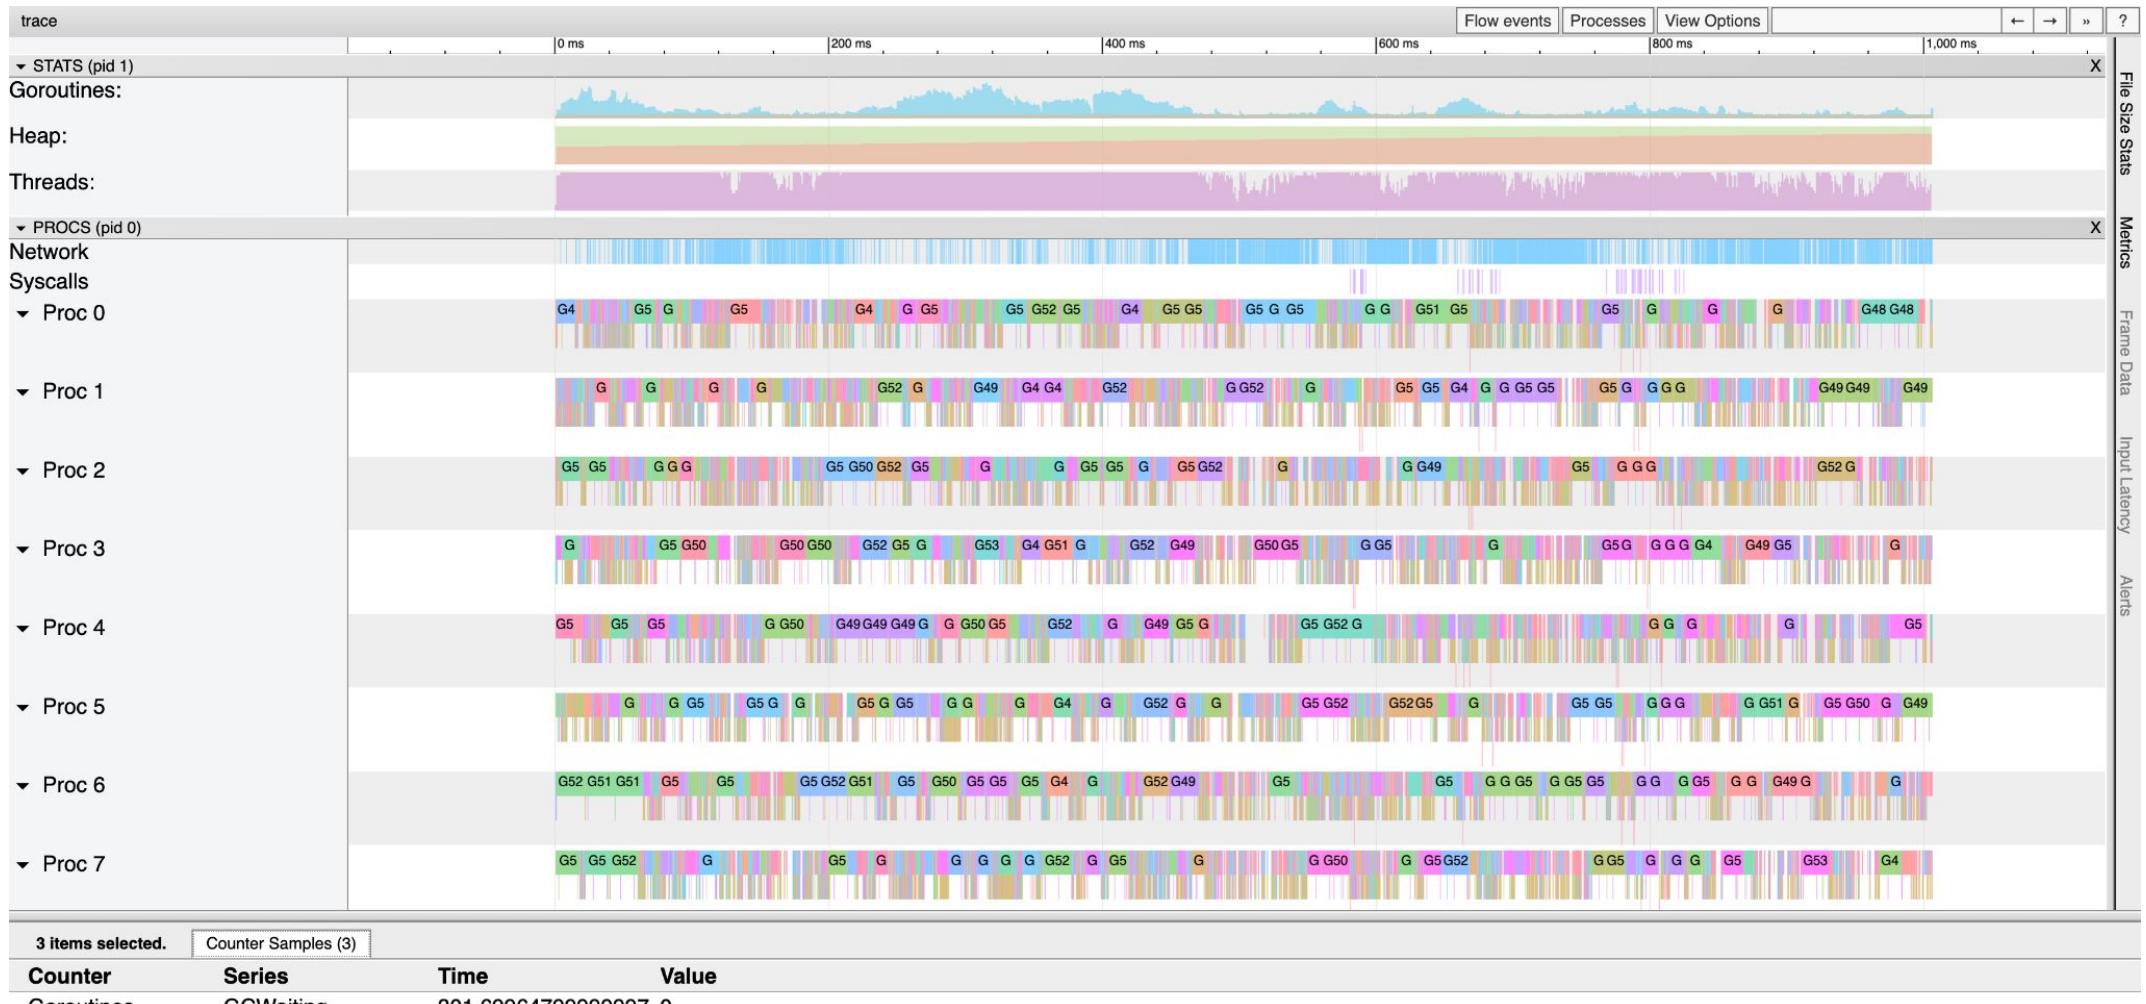
# runtime/trace has it all ...



# Text Analysis

```
$ go tool trace -d=parsed go.trace
...
M=1 P=1 G=2 StateTransition Time=100 GoID=2 Running->Waiting Reason="chan receive"
TransitionStack=
    runtime.chanrecv1 @ 0x1044f4873
        /code/github.com/golang/go/src/runtime/chan.go:509
testing.(*T).Run @ 0x1045f0ee7
    /code/github.com/golang/go/src/testing/testing.go:2005
...
M=1 P=1 G=1 StateTransition Time=300 GoID=2 Waiting->Runnable Reason=""
Stack=
    runtime.chansend1 @ 0x1044f3a17
        /code/github.com/golang/go/src/runtime/chan.go:161
testing.tRunner.func1.1 @ 0x1045f0b3f
    /code/github.com/golang/go/src/testing/testing.go:1849
M=1 P=0 G=-1 StateTransition Time=310 GoID=2 Runnable->Running Reason=""
```

# UI Analysis

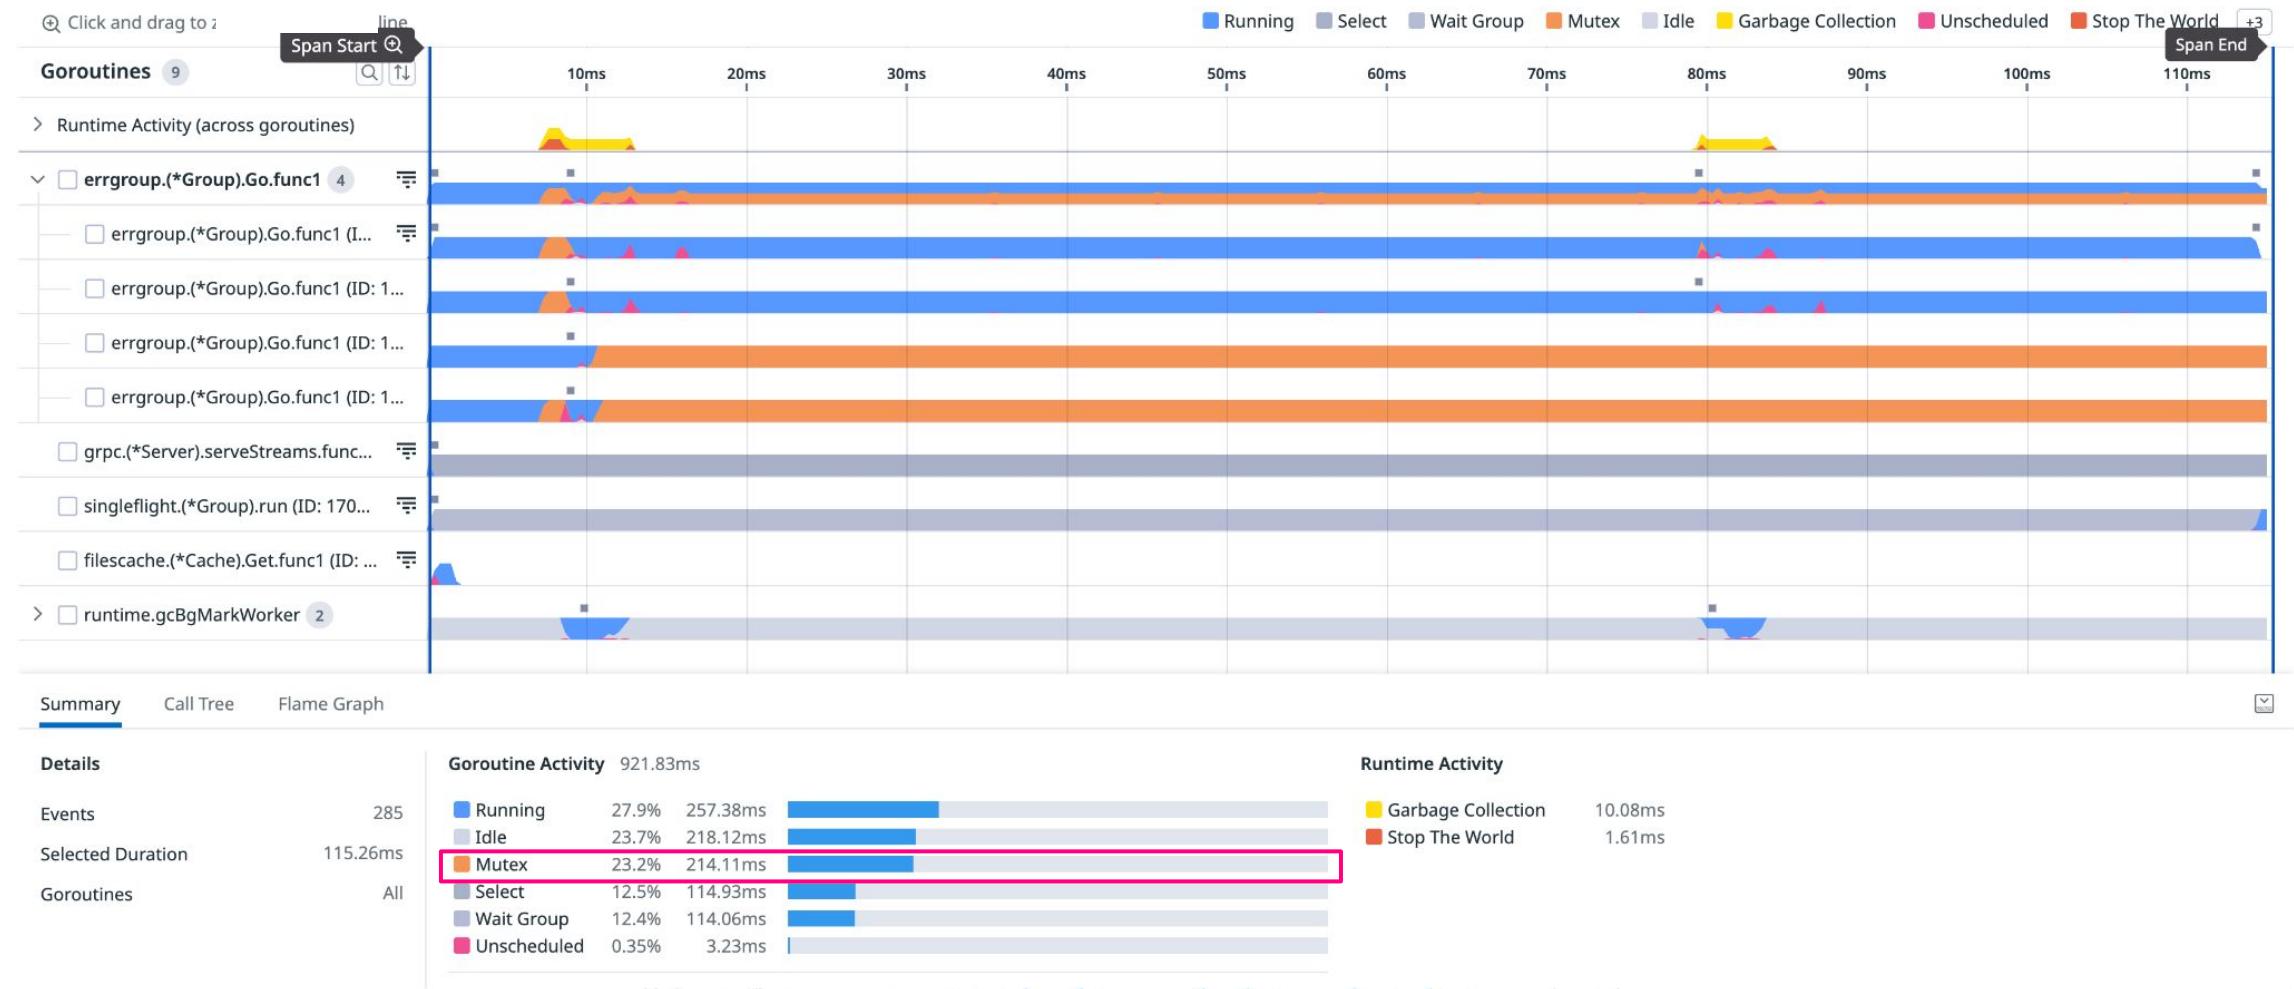


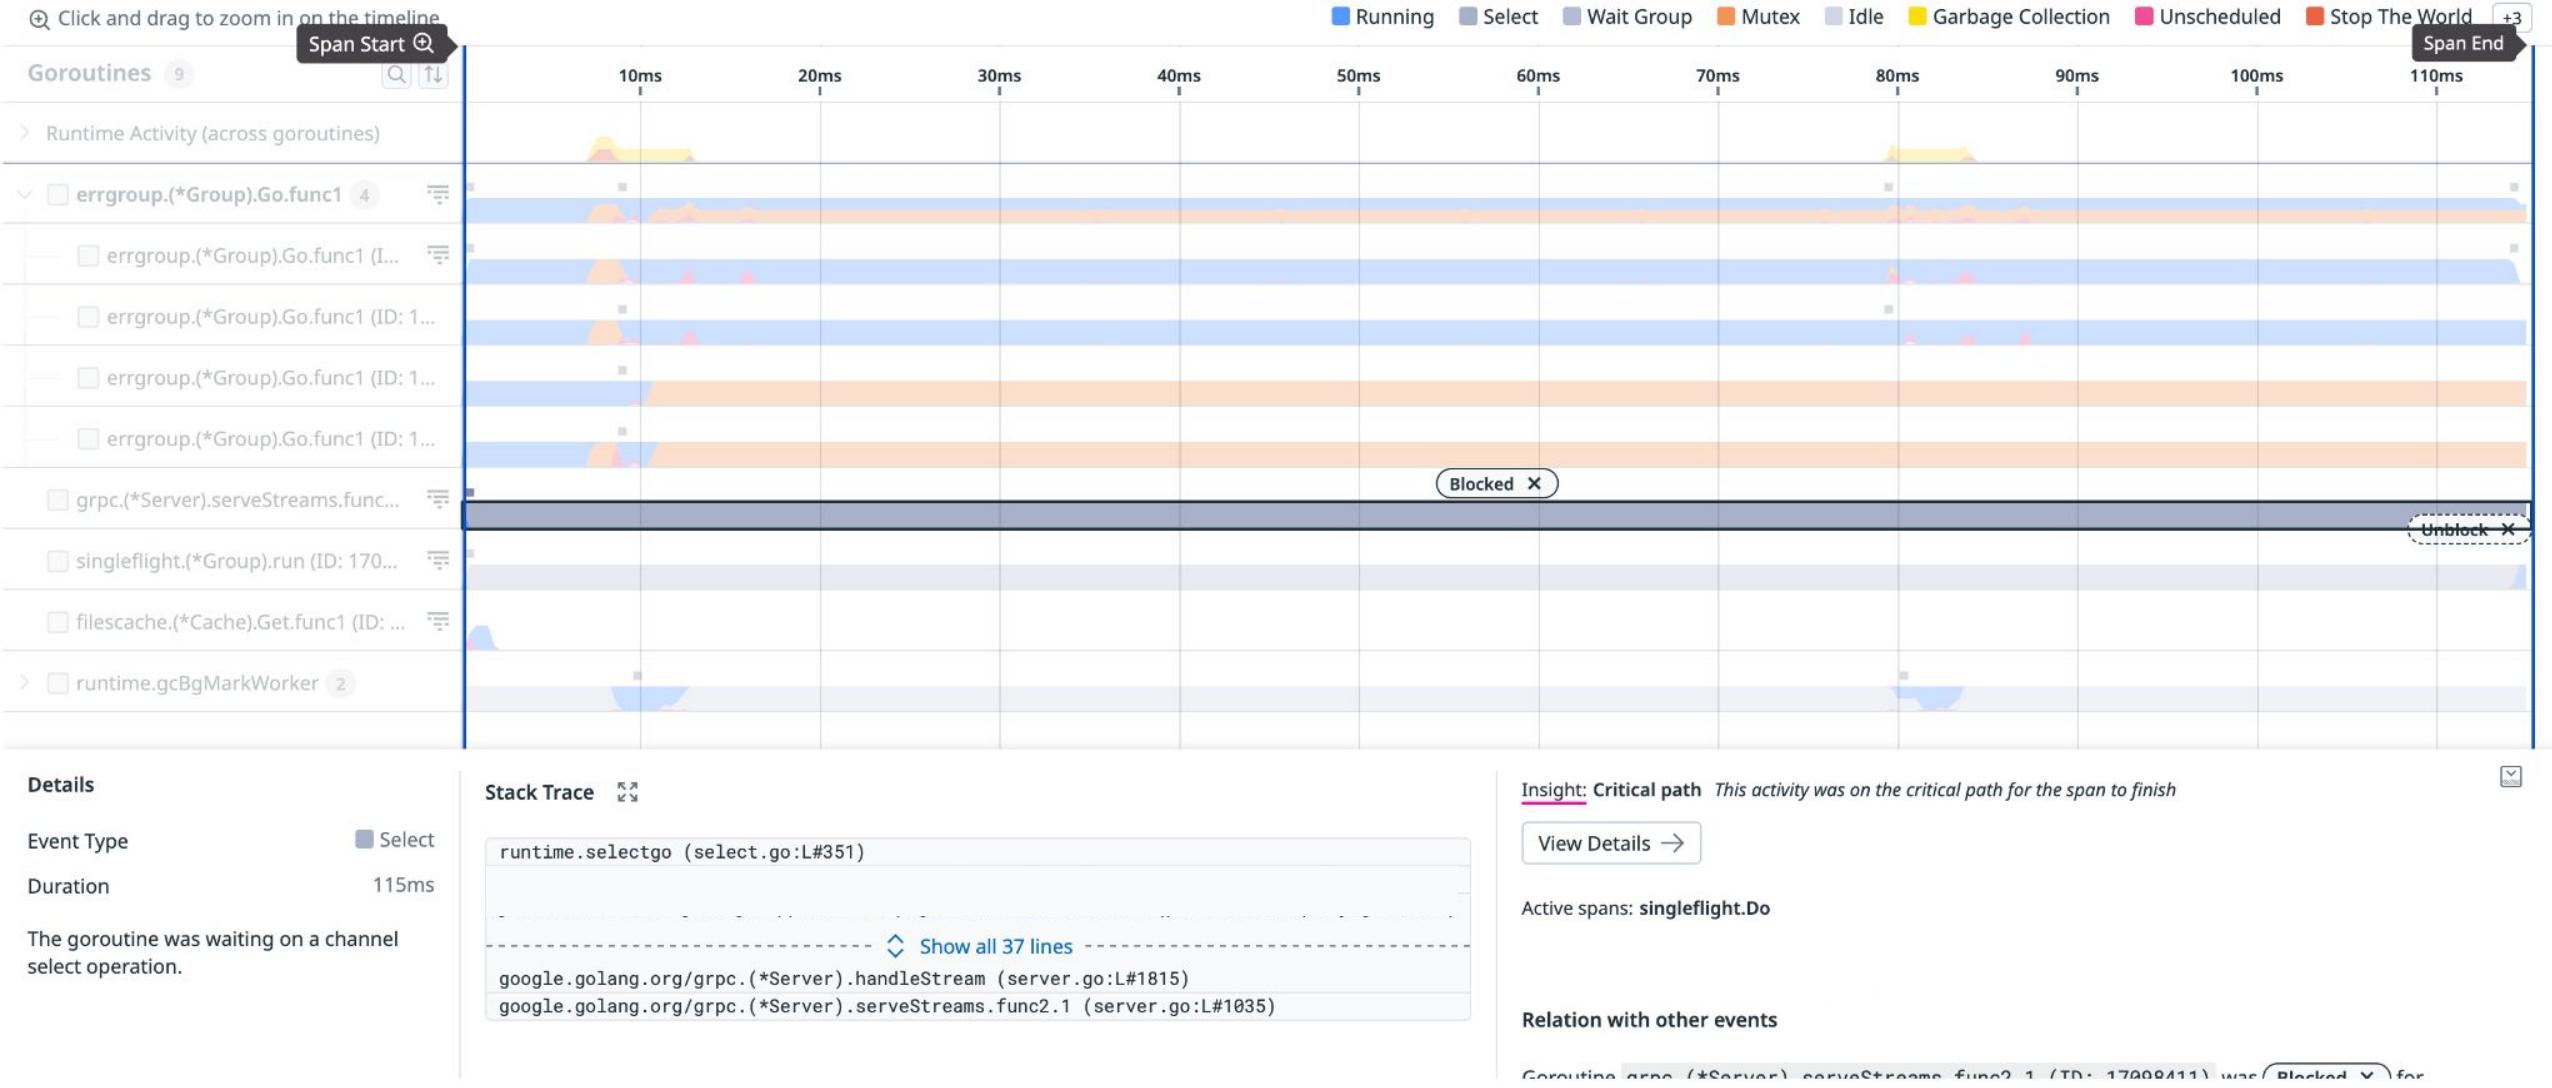
# Request Annotation

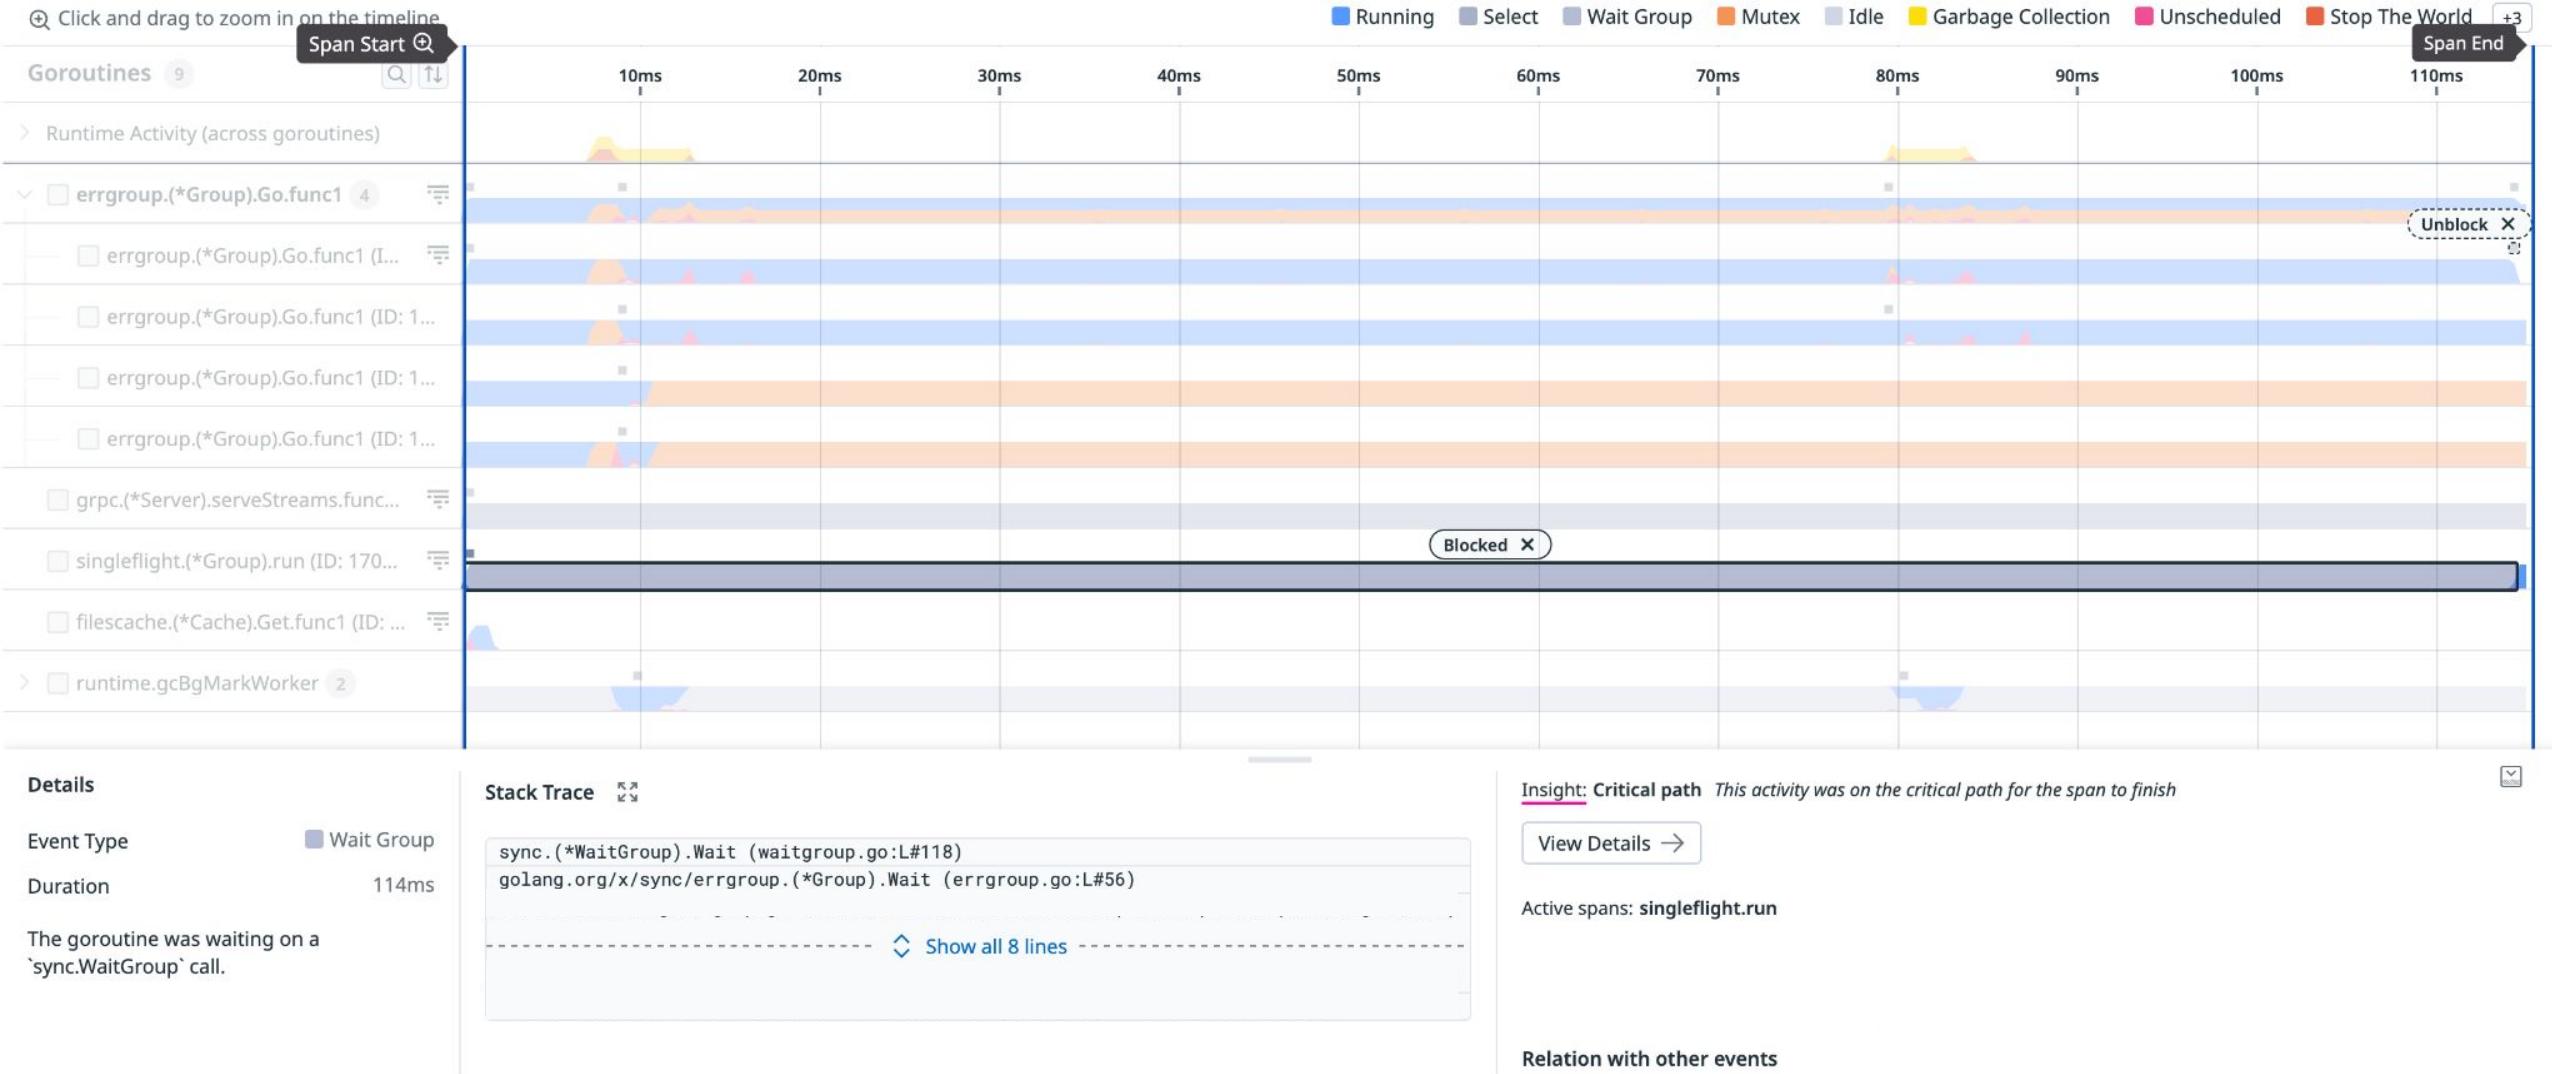
```
import (
    "net/http"
    "runtime/trace"
)

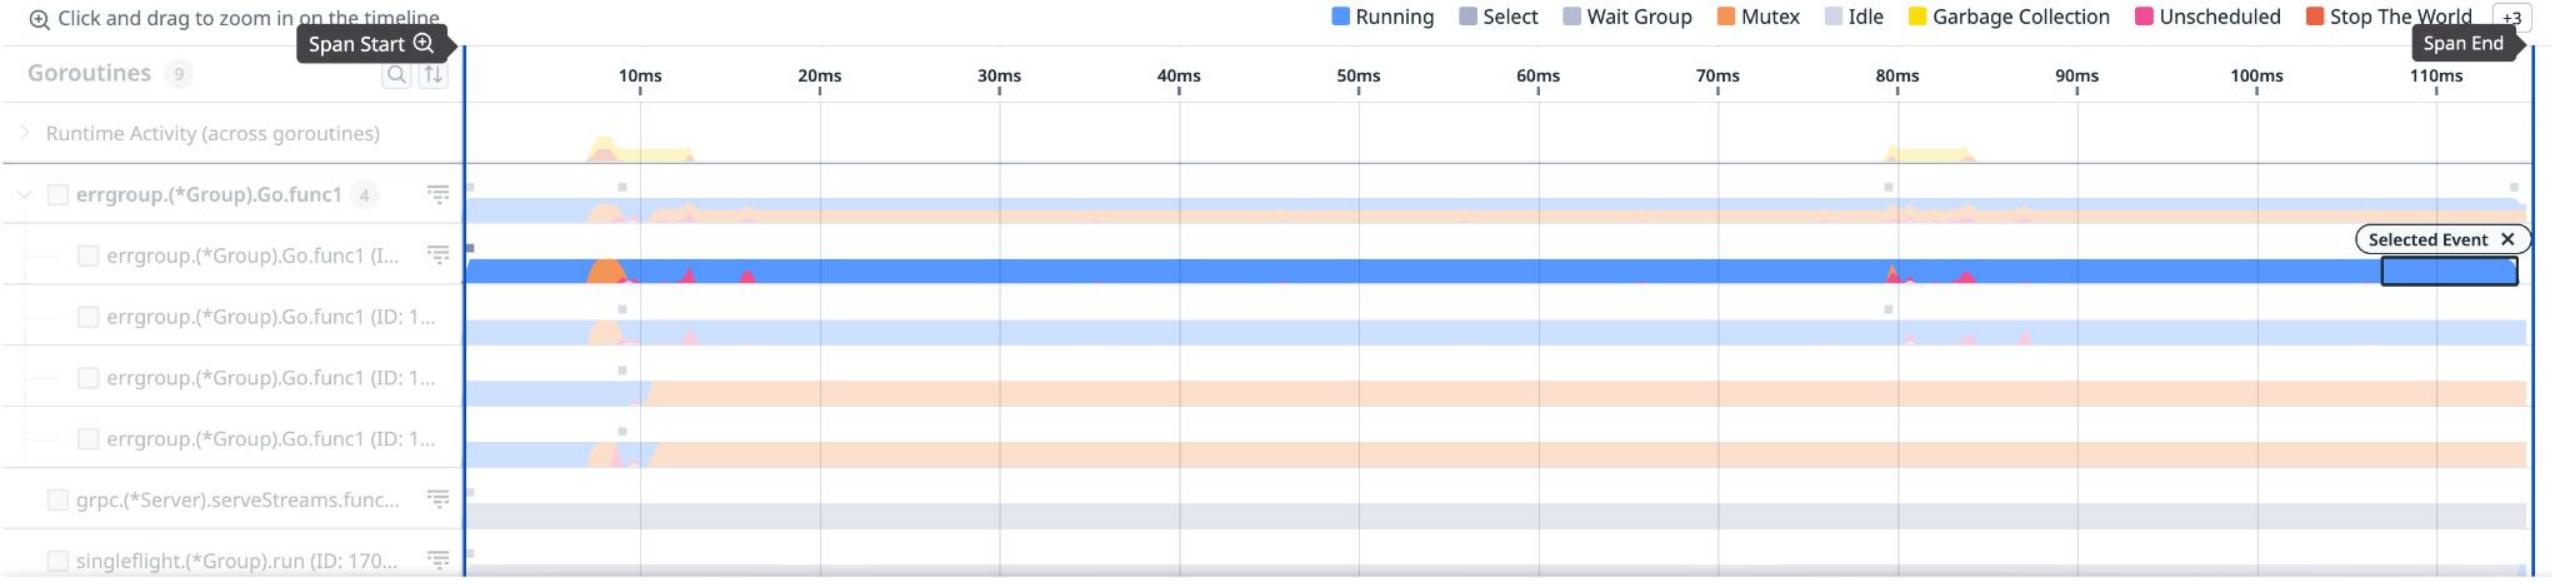
func TaskMiddleware(next http.HandlerFunc) http.HandlerFunc {
    return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
        ctx, task := trace.NewTask(r.Context(), "http.request")
        defer task.End()
        trace.Log(ctx, "http.request.url", r.URL.String())
        r = r.WithContext(ctx)
        next.ServeHTTP(w, r)
    })
}
```

# Datadog Timeline (Goroutine Oriented Trace Viewer)









## Details

Event Type

 Running

Duration

7ms

Cpu Sample

no

The goroutine was running on an OS thread. The stack trace was captured at the end of the event.

Stack Trace

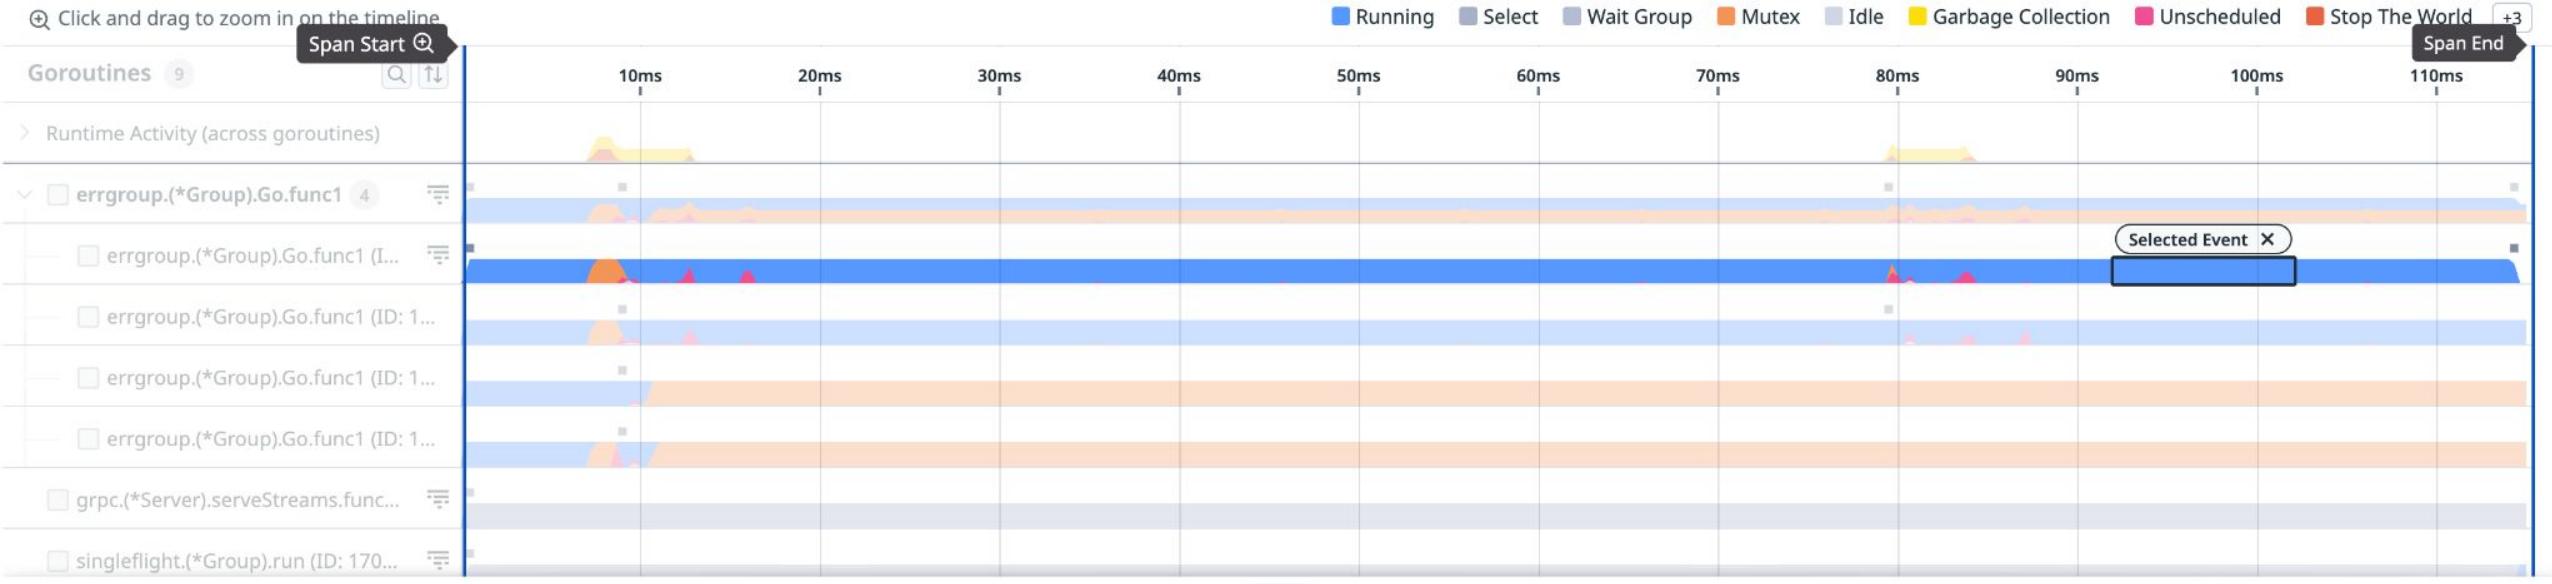
No Stack Trace available

Insight: Critical path This activity was on the critical path for the span to finish

[View Details →](#)

Active spans: singleflight.run

singleflight.Do



## Details

Event Type

Duration 10ms

Cpu Sample yes

The goroutine was running on an OS thread. The stack trace was captured at the end of the event.

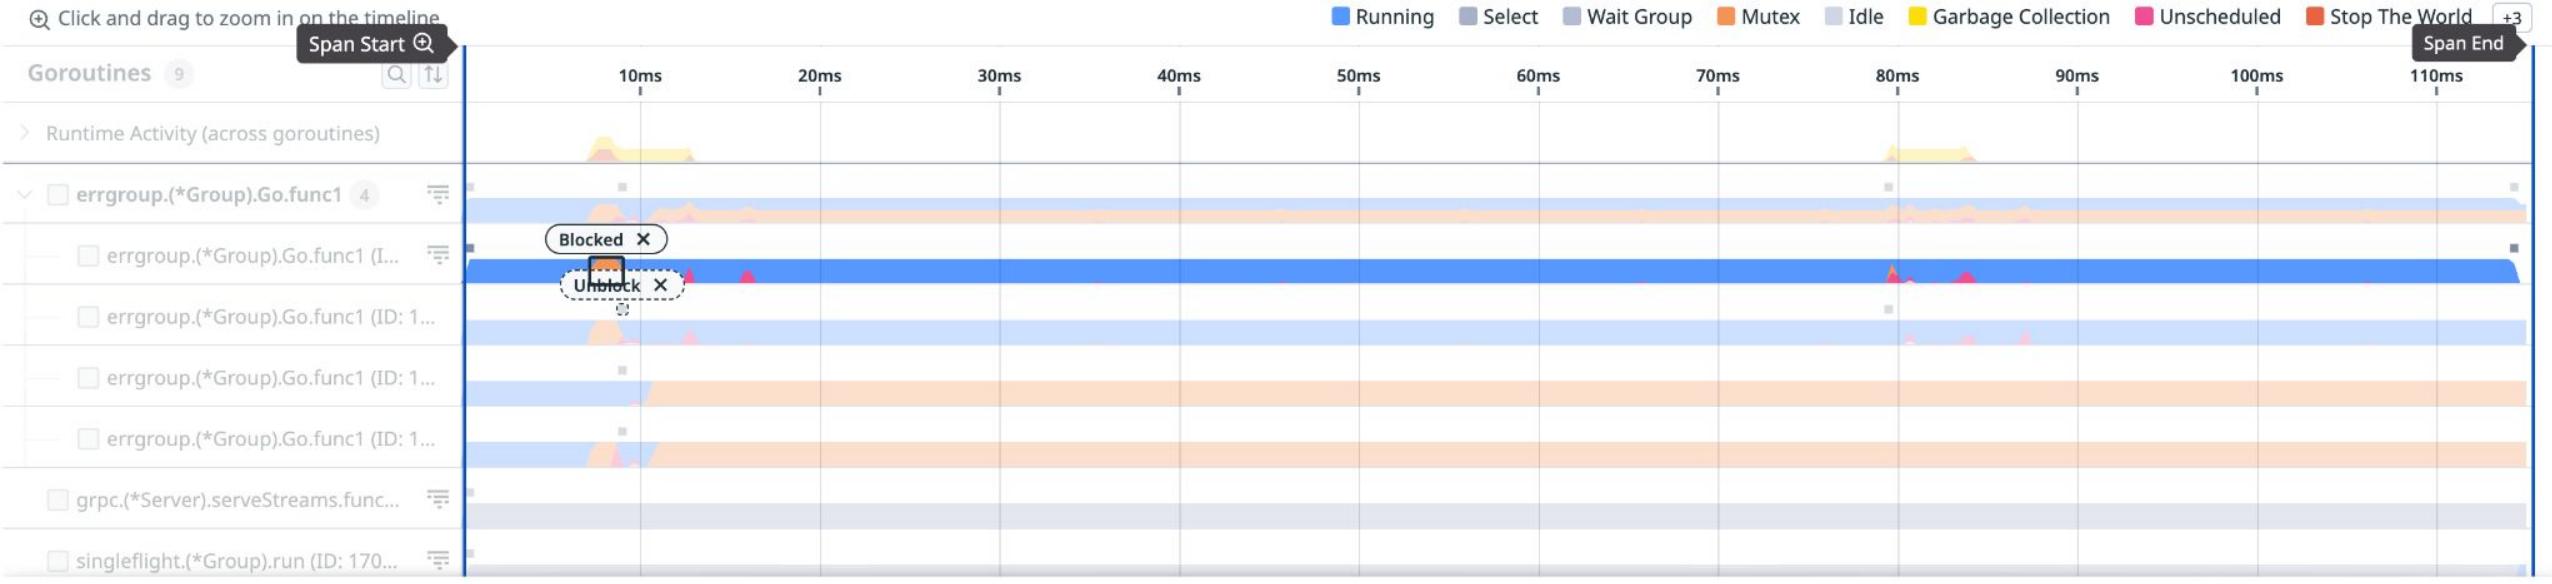
## Stack Trace

```
runtime.unlock2 (lock_spinbit.go:L#262)
runtime.unlockWithRank (lockrank_off.go:L#35)
runtime.unlock (lock_spinbit.go:L#253)
----- ▾ Show all 21 lines ▾ -----
golang.org/x/sync/errgroup.(*Group).Go.func1 (errgroup.go:L#93)
```

**Insight: Critical path** This activity was on the critical path for the span to finish

[View Details →](#)

Active spans: **singleflight.run**  
**singleflight.Do**



## Details

Event Type Mutex

Duration 2ms

The goroutine was blocked while trying to acquire a lock that was held by another goroutine.

## Stack Trace

```
runtime.gcStart (mgc.go:L#668)
runtime.mallocgcSmallScanNoHeader (malloc.go:L#1425)
runtime.mallocgc (malloc.go:L#1058)

golang.org/x/sync/errgroup.(*Group).Go.func1 (errgroup.go:L#93)
```

Show all 19 lines

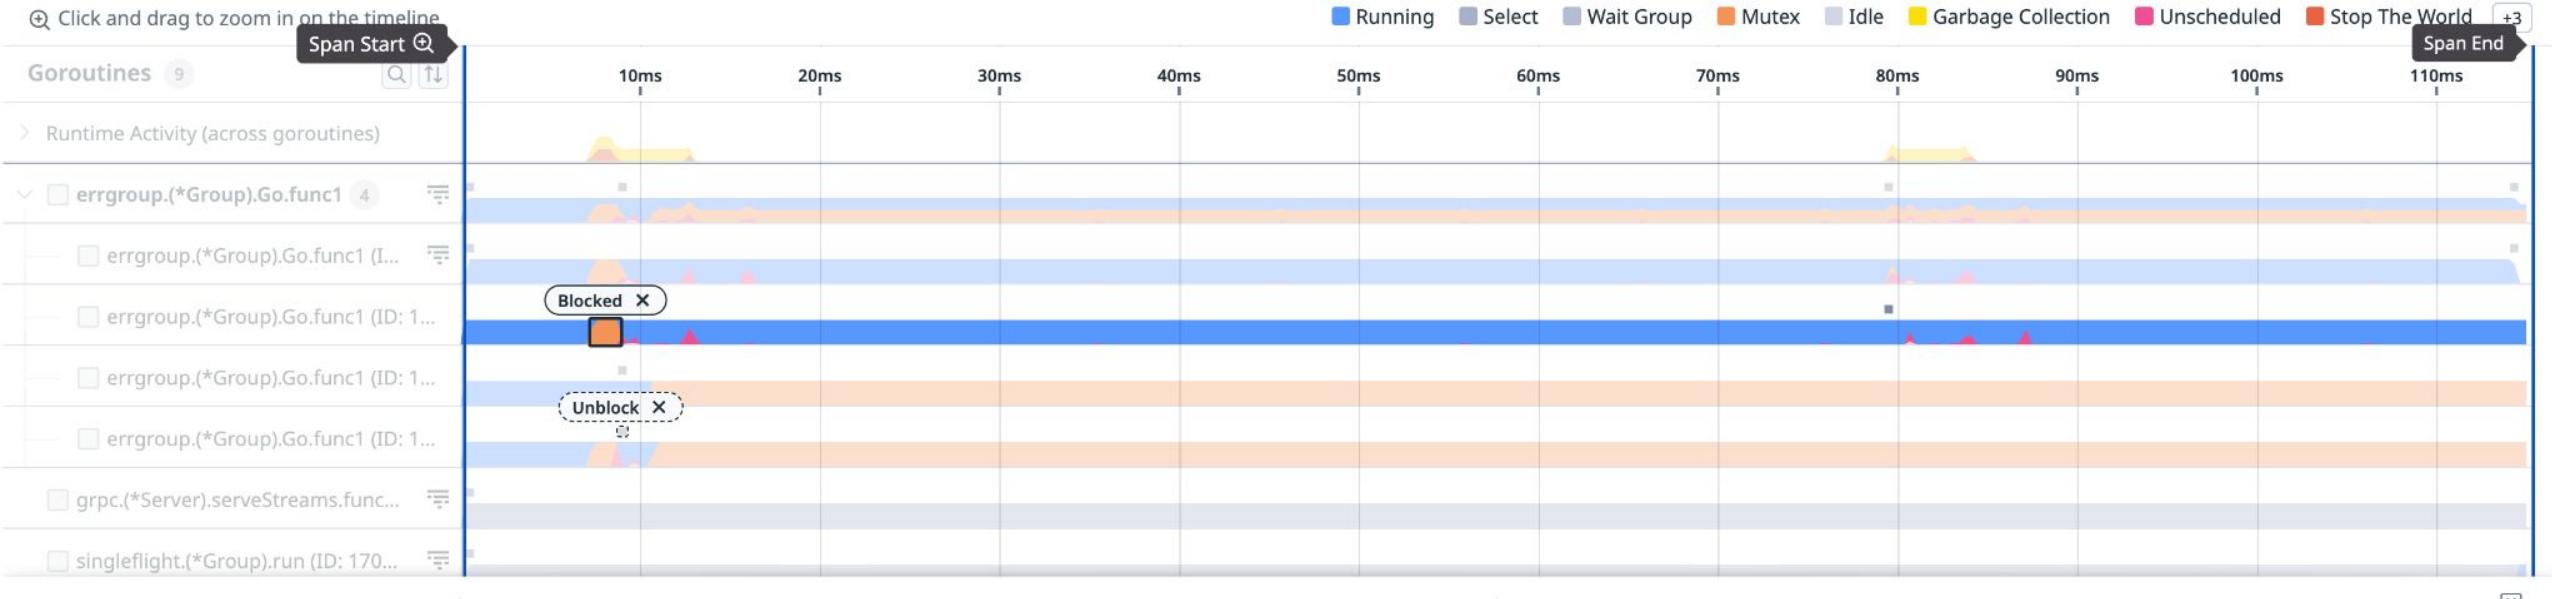
Insight: Critical path This activity was on the critical path for the span to finish

[View Details →](#)

Active spans: singleflight.run  
singleflight.Do

## Relation with other events

Goroutine errgroup.(\*Group).Go.func1 (ID: 17098413) was Blocked for 1.7ms and was then Unblocked by goroutine errgroup.(\*Group).Go.func1 (ID: 17099058)



## Details

Event Type

Duration 2ms

The goroutine was blocked while trying to acquire a lock that was held by another goroutine.

## Stack Trace

```
runtime.gcStart (mgc.go:L#668)
runtime.mallocgcSmallScanNoHeader (malloc.go:L#1425)
runtime.mallocgc (malloc.go:L#1058)

----- Show all 19 lines -----
```

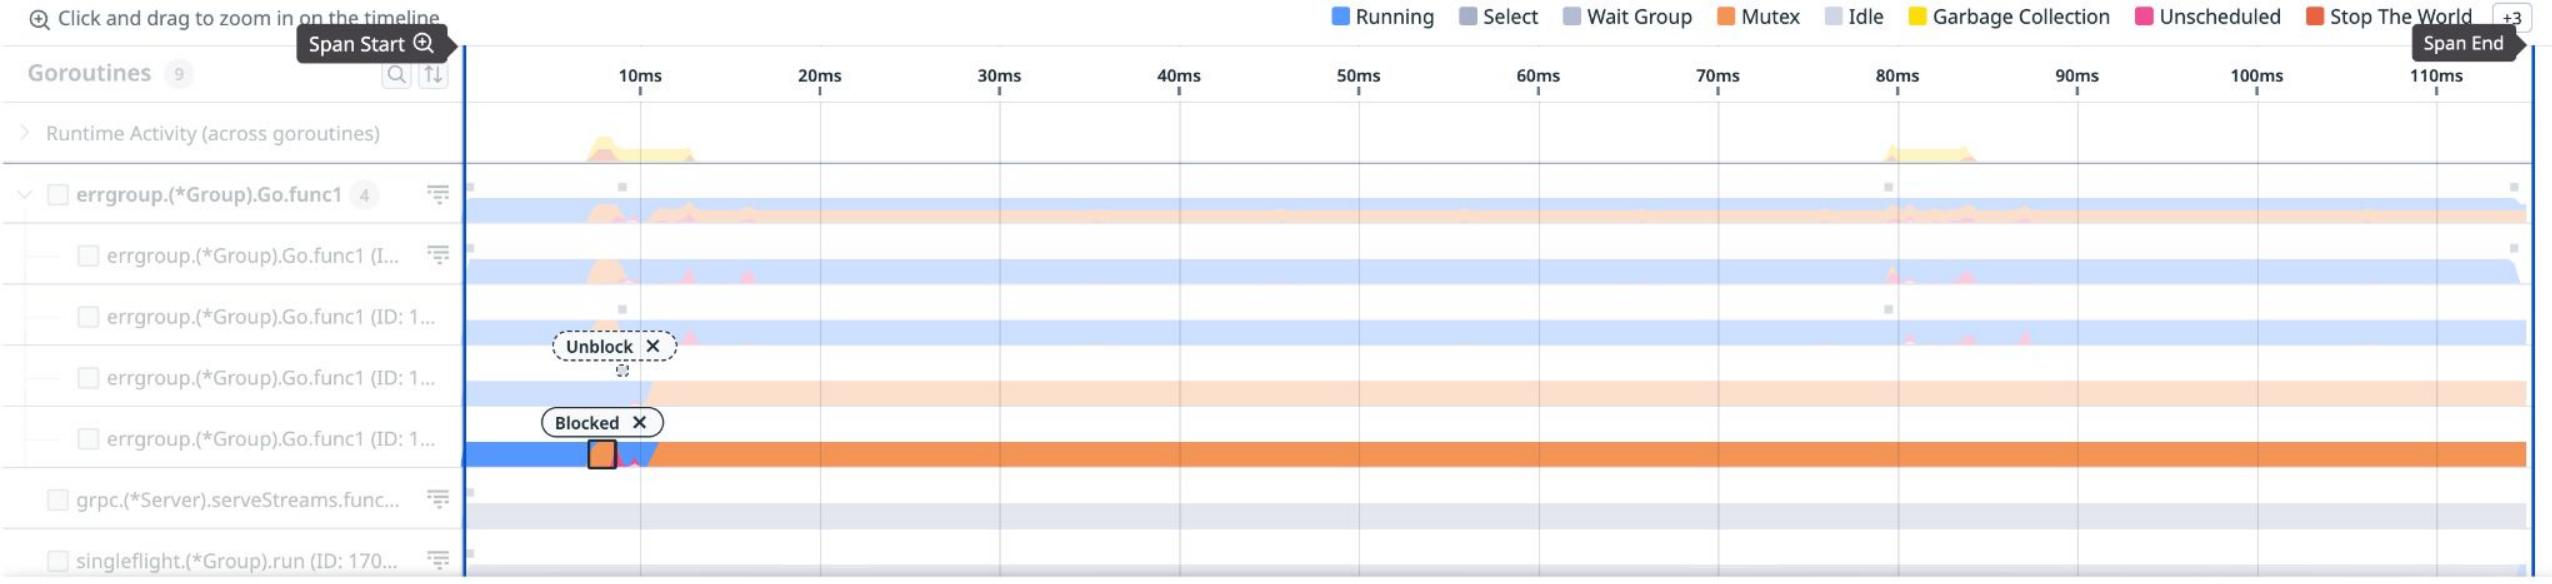
golang.org/x/sync/errgroup.(\*Group).Go.func1 (errgroup.go:L#93)

**Insight: Critical path** This activity was on the critical path for the span to finish

[View Details →](#)

**Relation with other events**

Goroutine errgroup.(\*Group).Go.func1 (ID: 17099058) was **Blocked X** for 1.61ms and was then **Unblocked X** by goroutine errgroup.(\*Group).Go.func1 (ID: 17099049)



## Details

Event Type Mutex

Duration 1ms

The goroutine was blocked while trying to acquire a lock that was held by another goroutine.

## Stack Trace

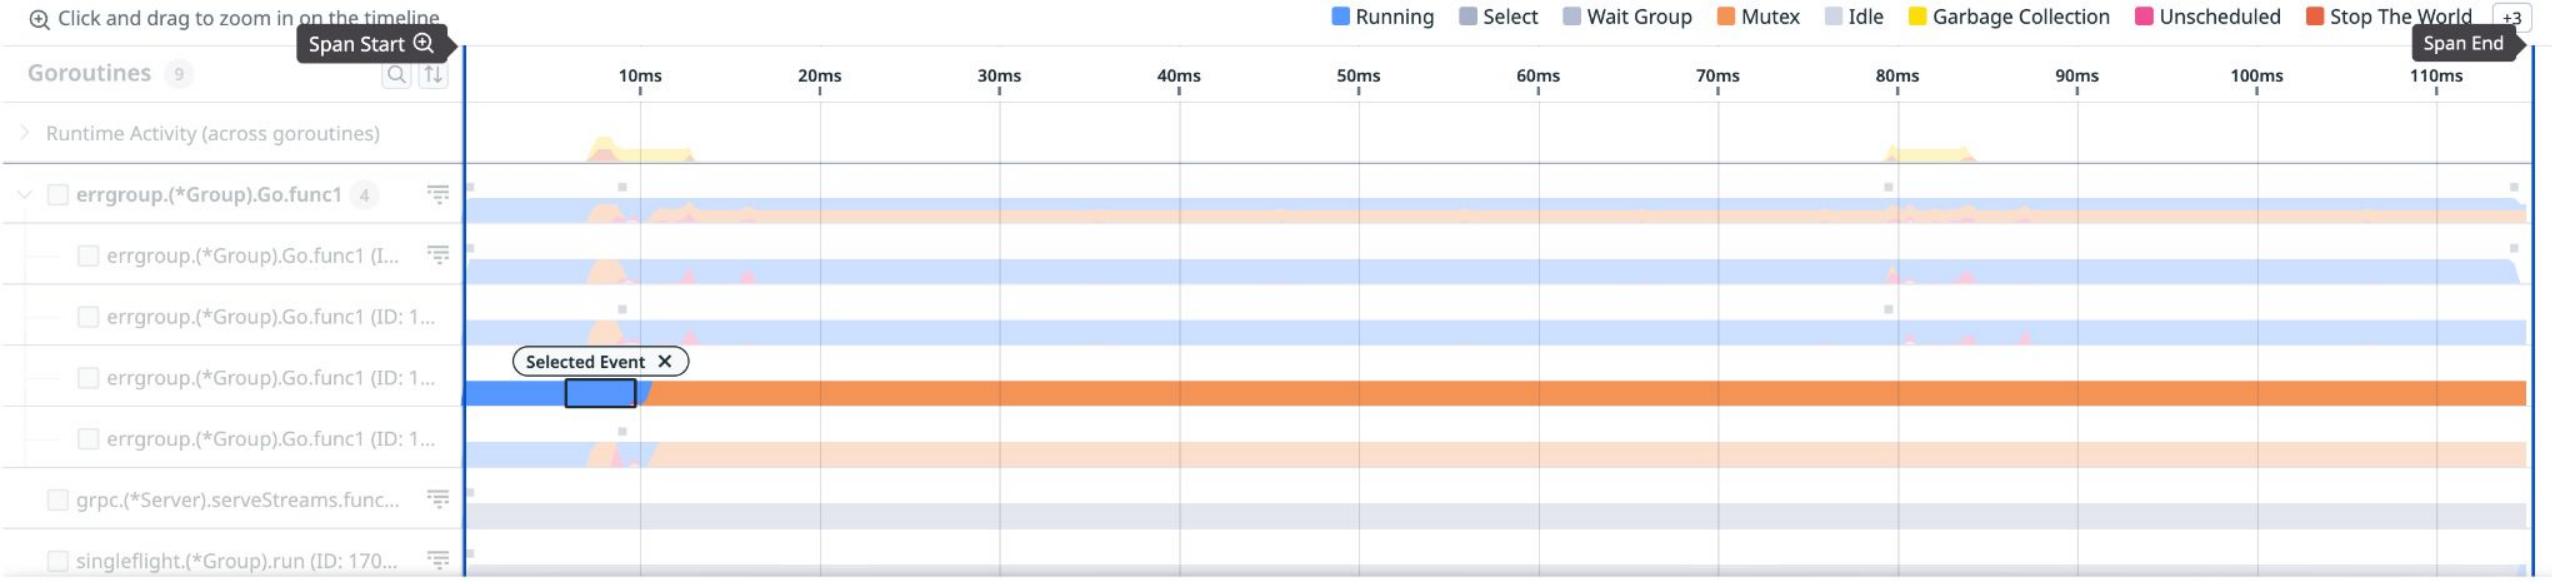
```
runtime.gcStart (mgc.go:L#668)
runtime.mallocgcLarge (malloc.go:L#1591)
runtime.mallocgc (malloc.go:L#1063)
----- 
golang.org/x/sync/errgroup.(*Group).Go.func1 (errgroup.go:L#93)
```

Insight: Critical path This activity was on the critical path for the span to finish

[View Details →](#)

**Relation with other events**

Goroutine errgroup.(\*Group).Go.func1 (ID: 17099049) was Blocked for 1.31ms and was then Unblocked by goroutine errgroup.(\*Group).Go.func1 (ID: 17099047)



## Details

Event Type

 Running

Duration

4ms

Cpu Sample

no

The goroutine was running on an OS thread. The stack trace was captured at the end of the event.

Stack Trace

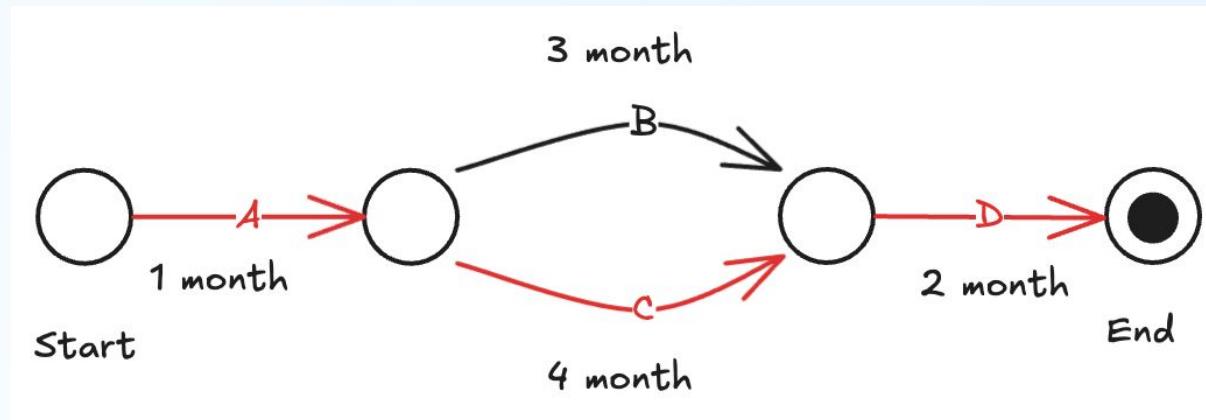
No Stack Trace available

Insight: Critical path This activity was on the critical path for the span to finish

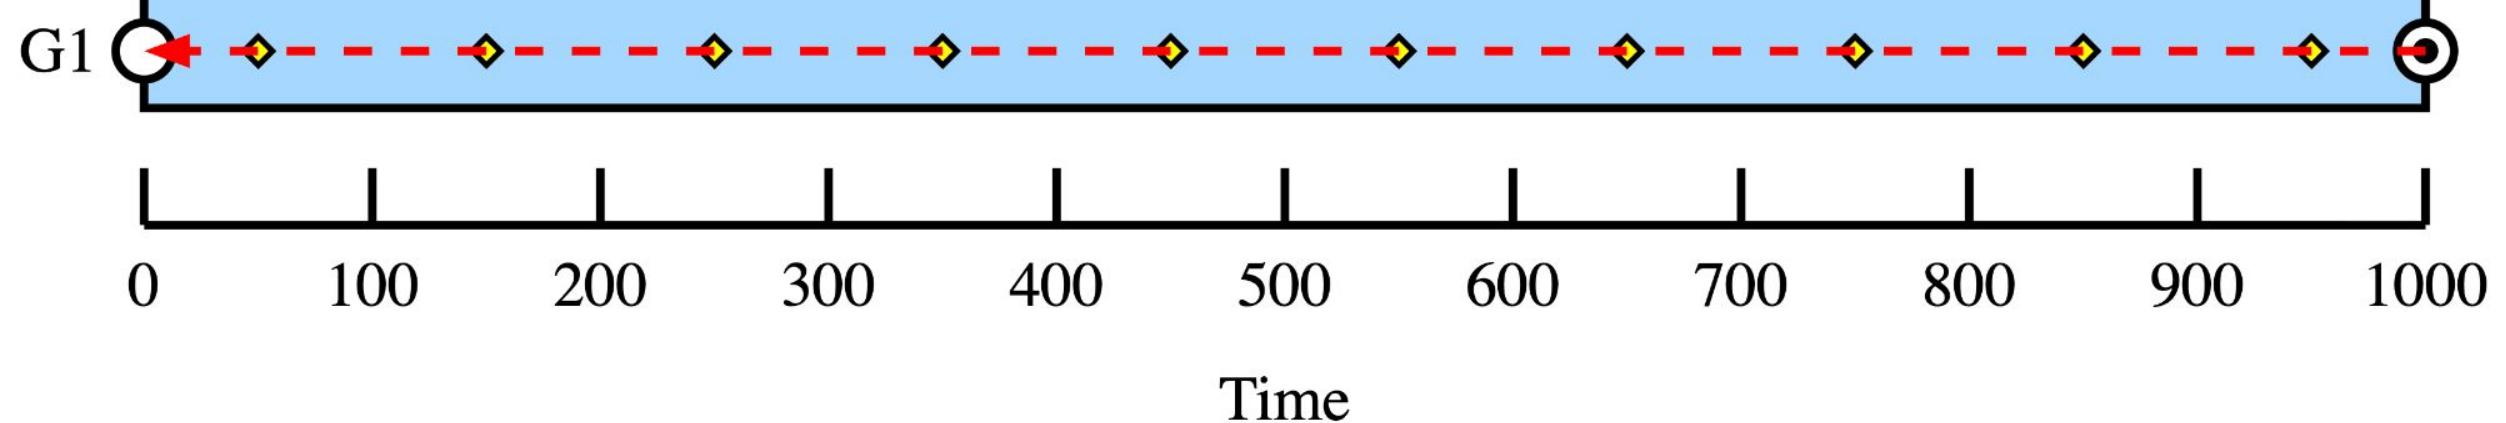
[View Details →](#)

# Critical Path Analysis

- The precursors of what came to be known as critical path were developed and put into practice by DuPont between 1940 and 1943 and contributed to the success of the Manhattan Project.
- A critical path is determined by identifying the longest stretch of dependent activities.

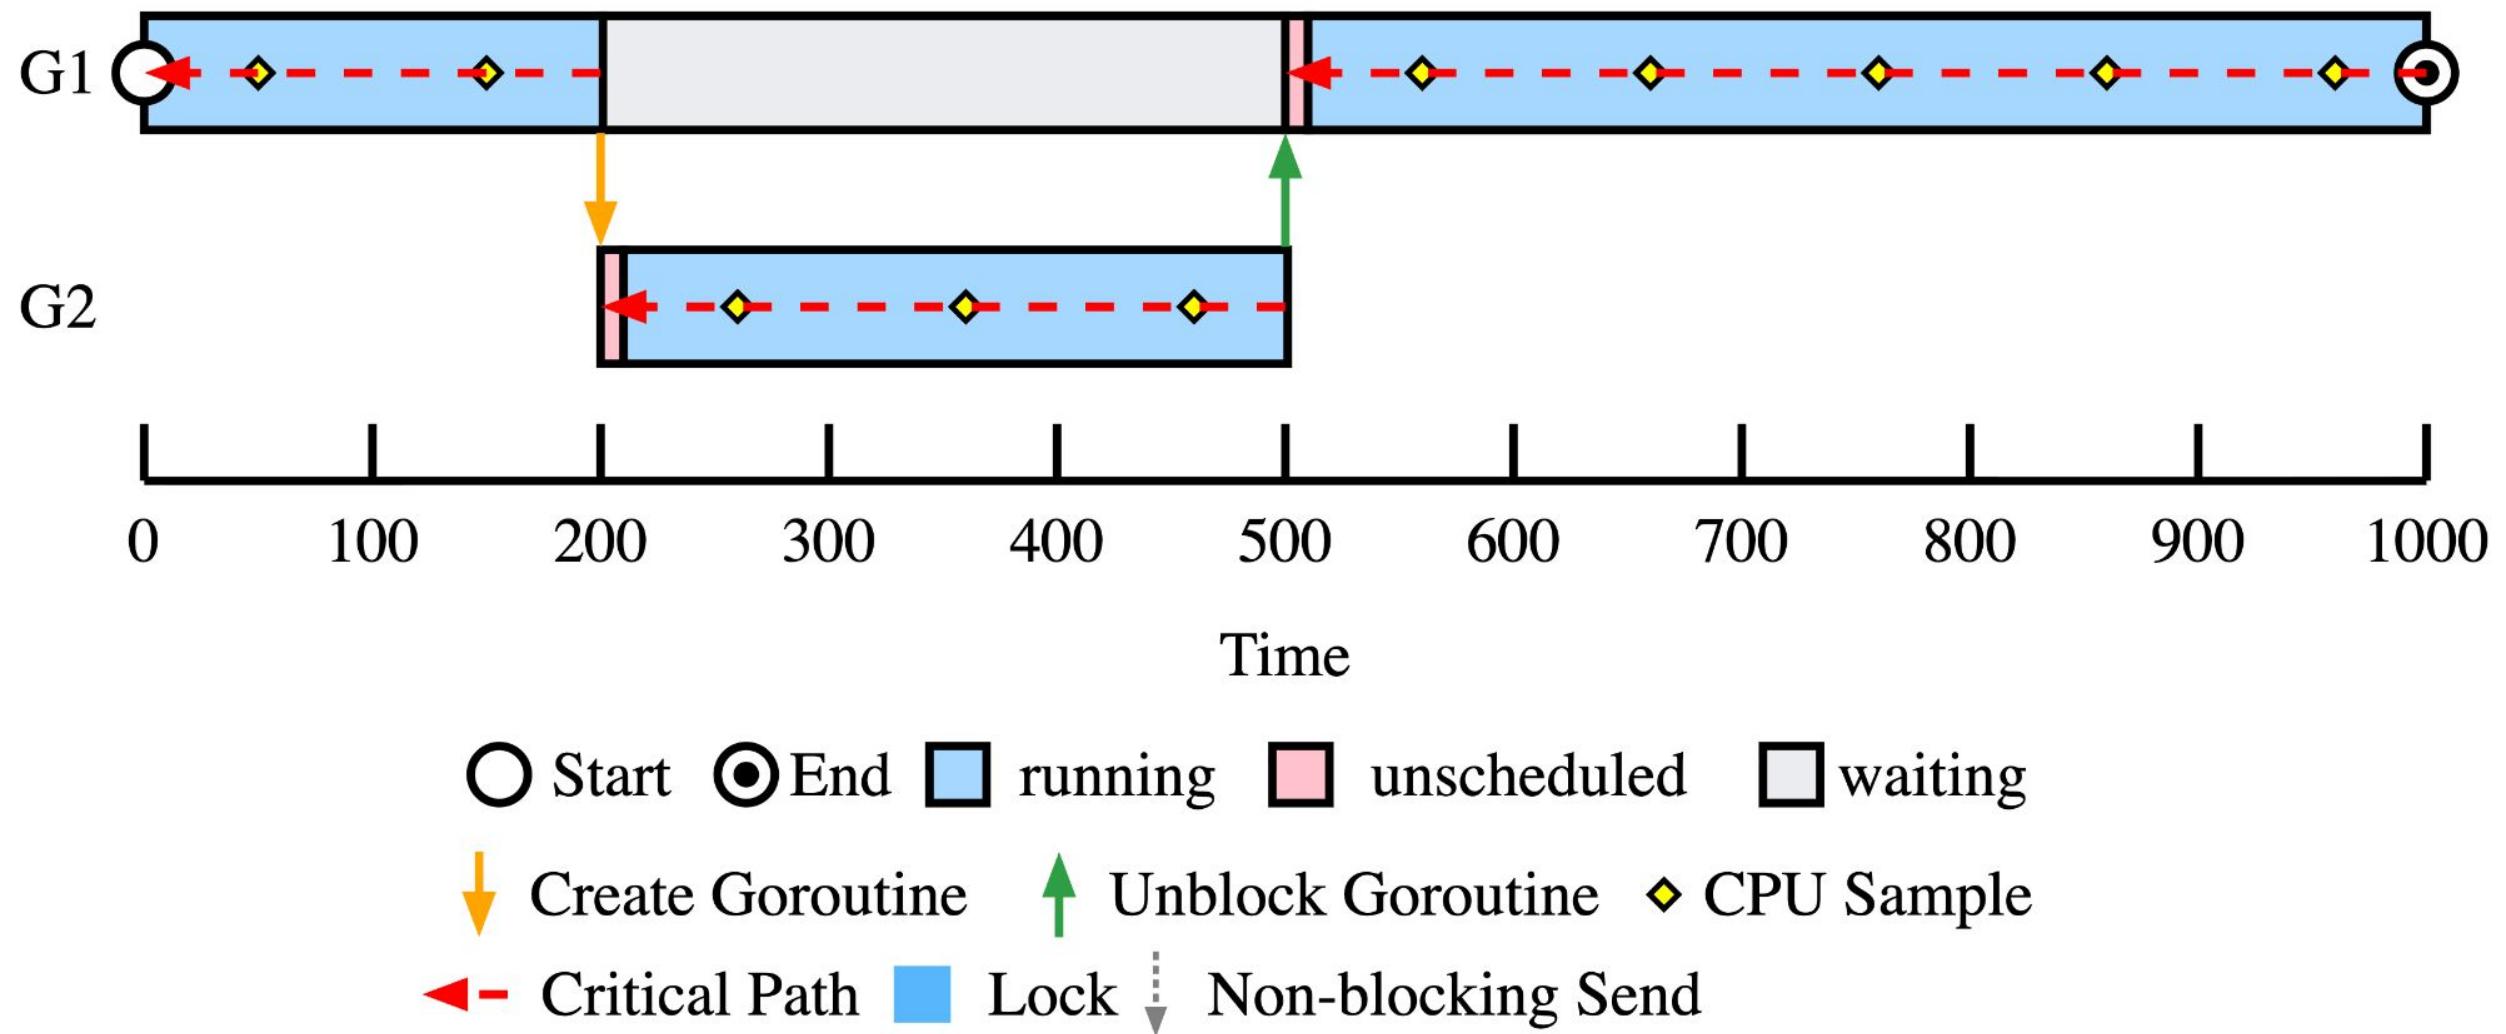


# OneGoroutine

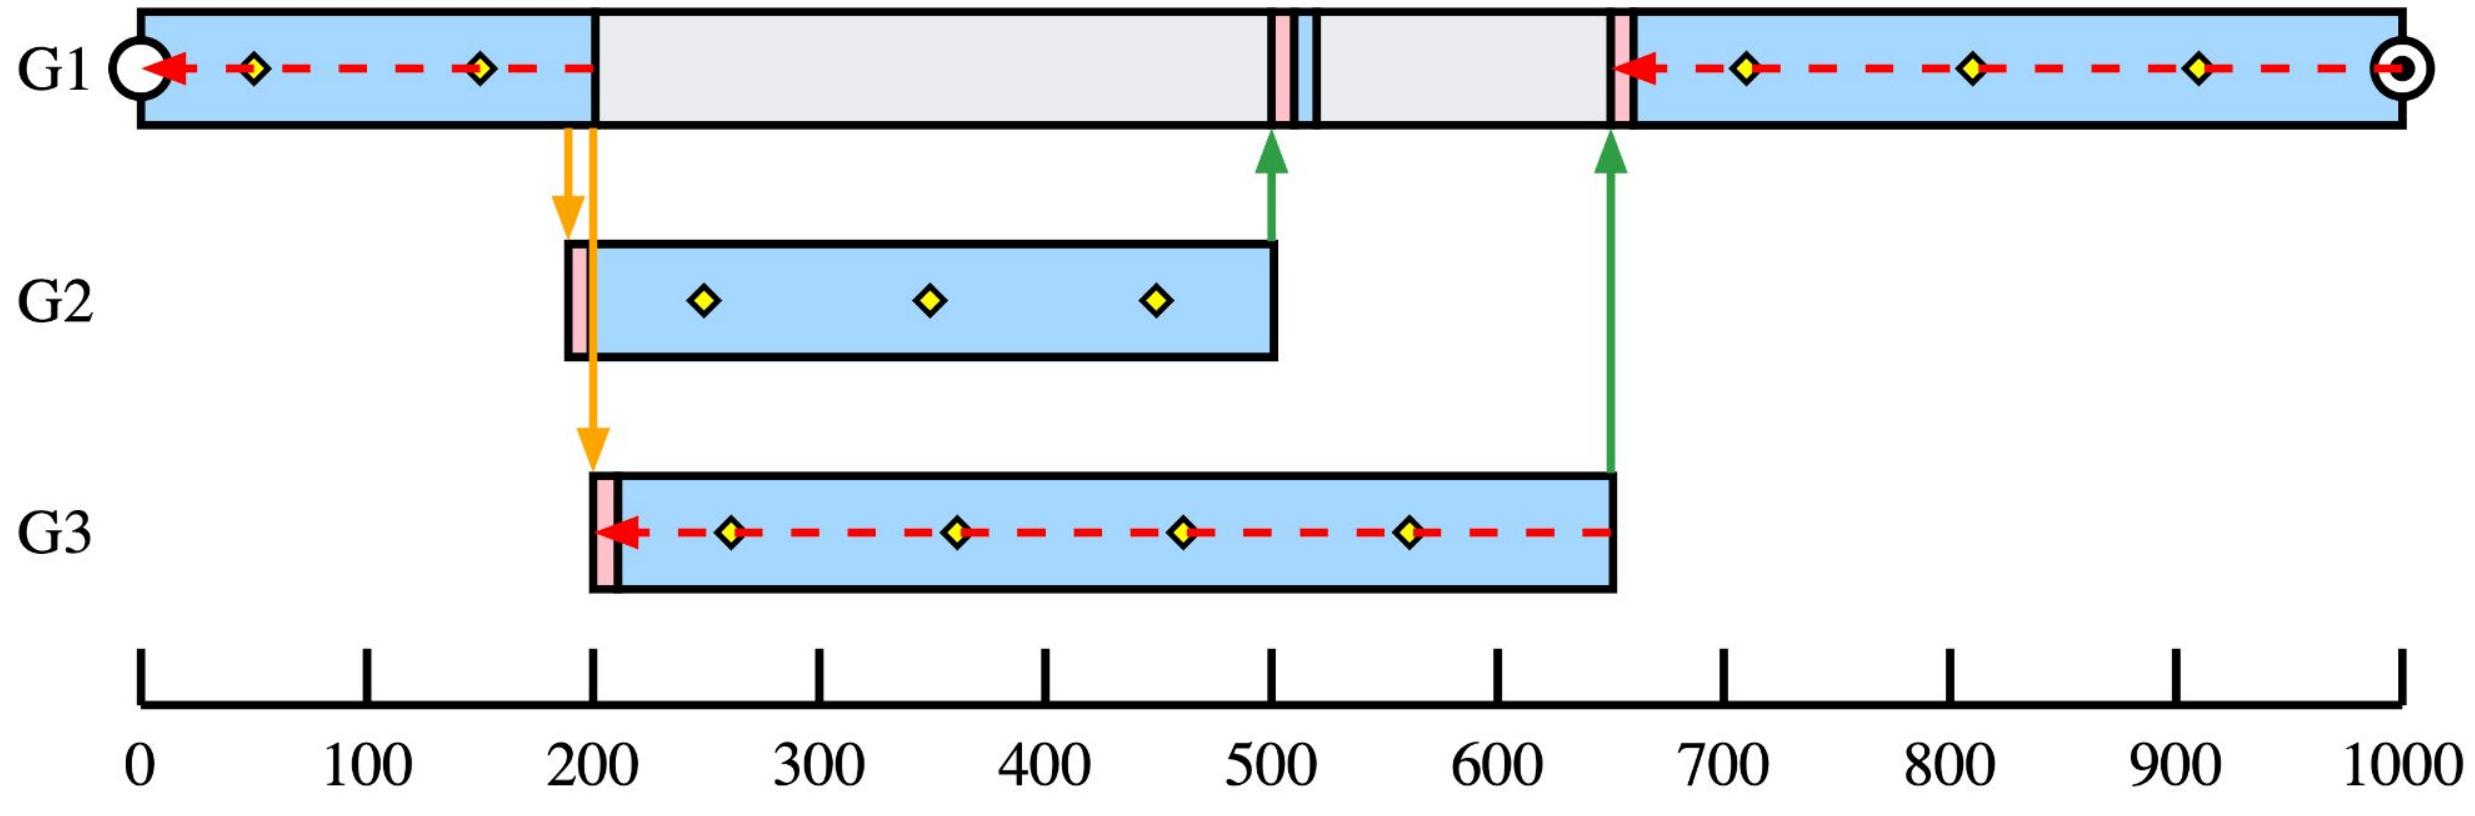


○ Start   ● End   ■ running   ■ unscheduled   ■ waiting  
↓ Create Goroutine   ↑ Unblock Goroutine   ◆ CPU Sample  
← Critical Path   ■ Lock   ↓ Non-blocking Send

# ChildGoroutine

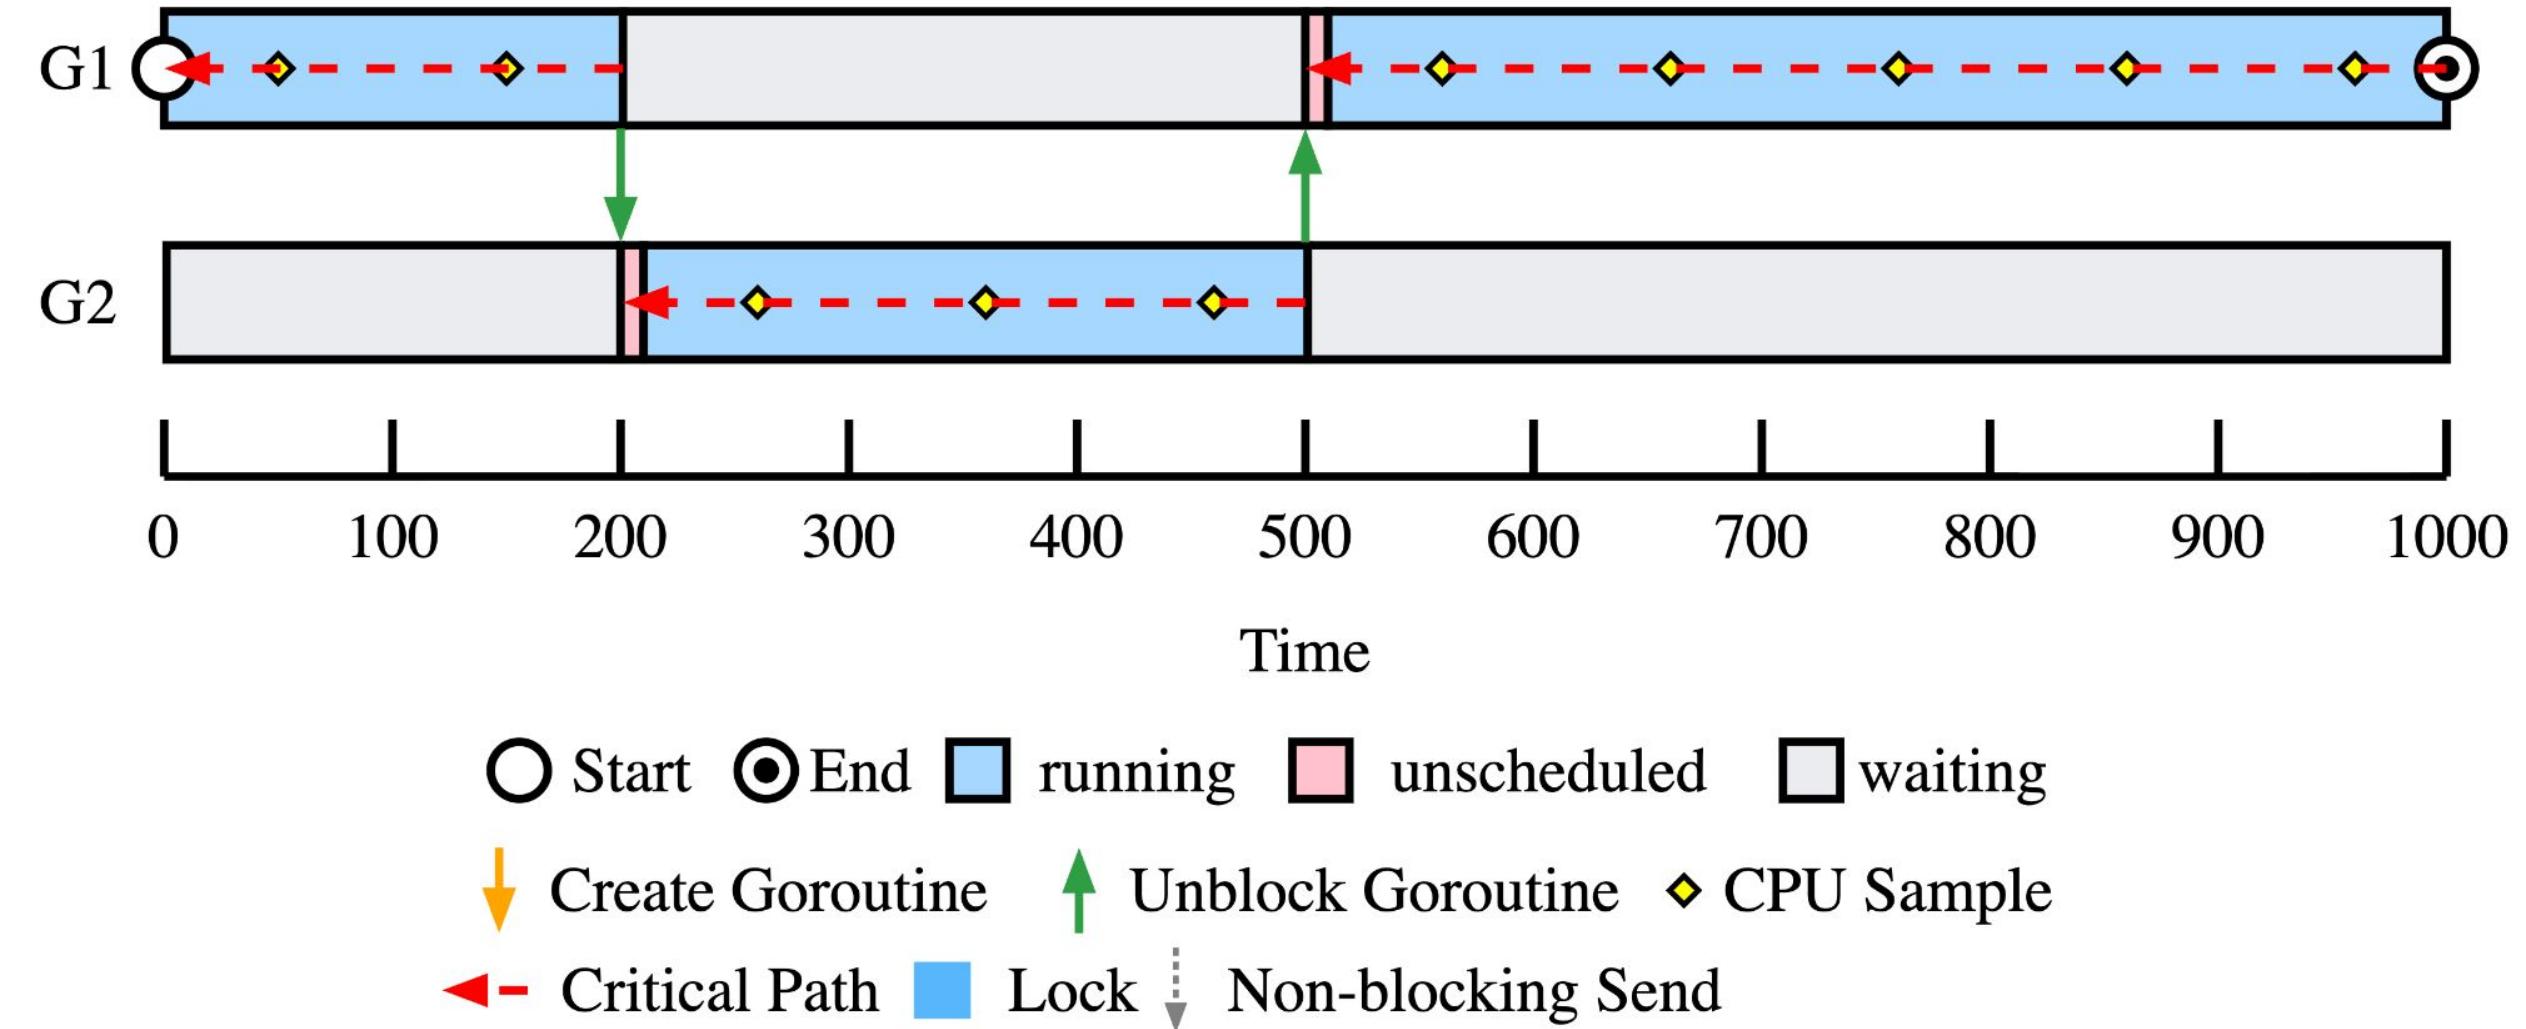


# FanOut

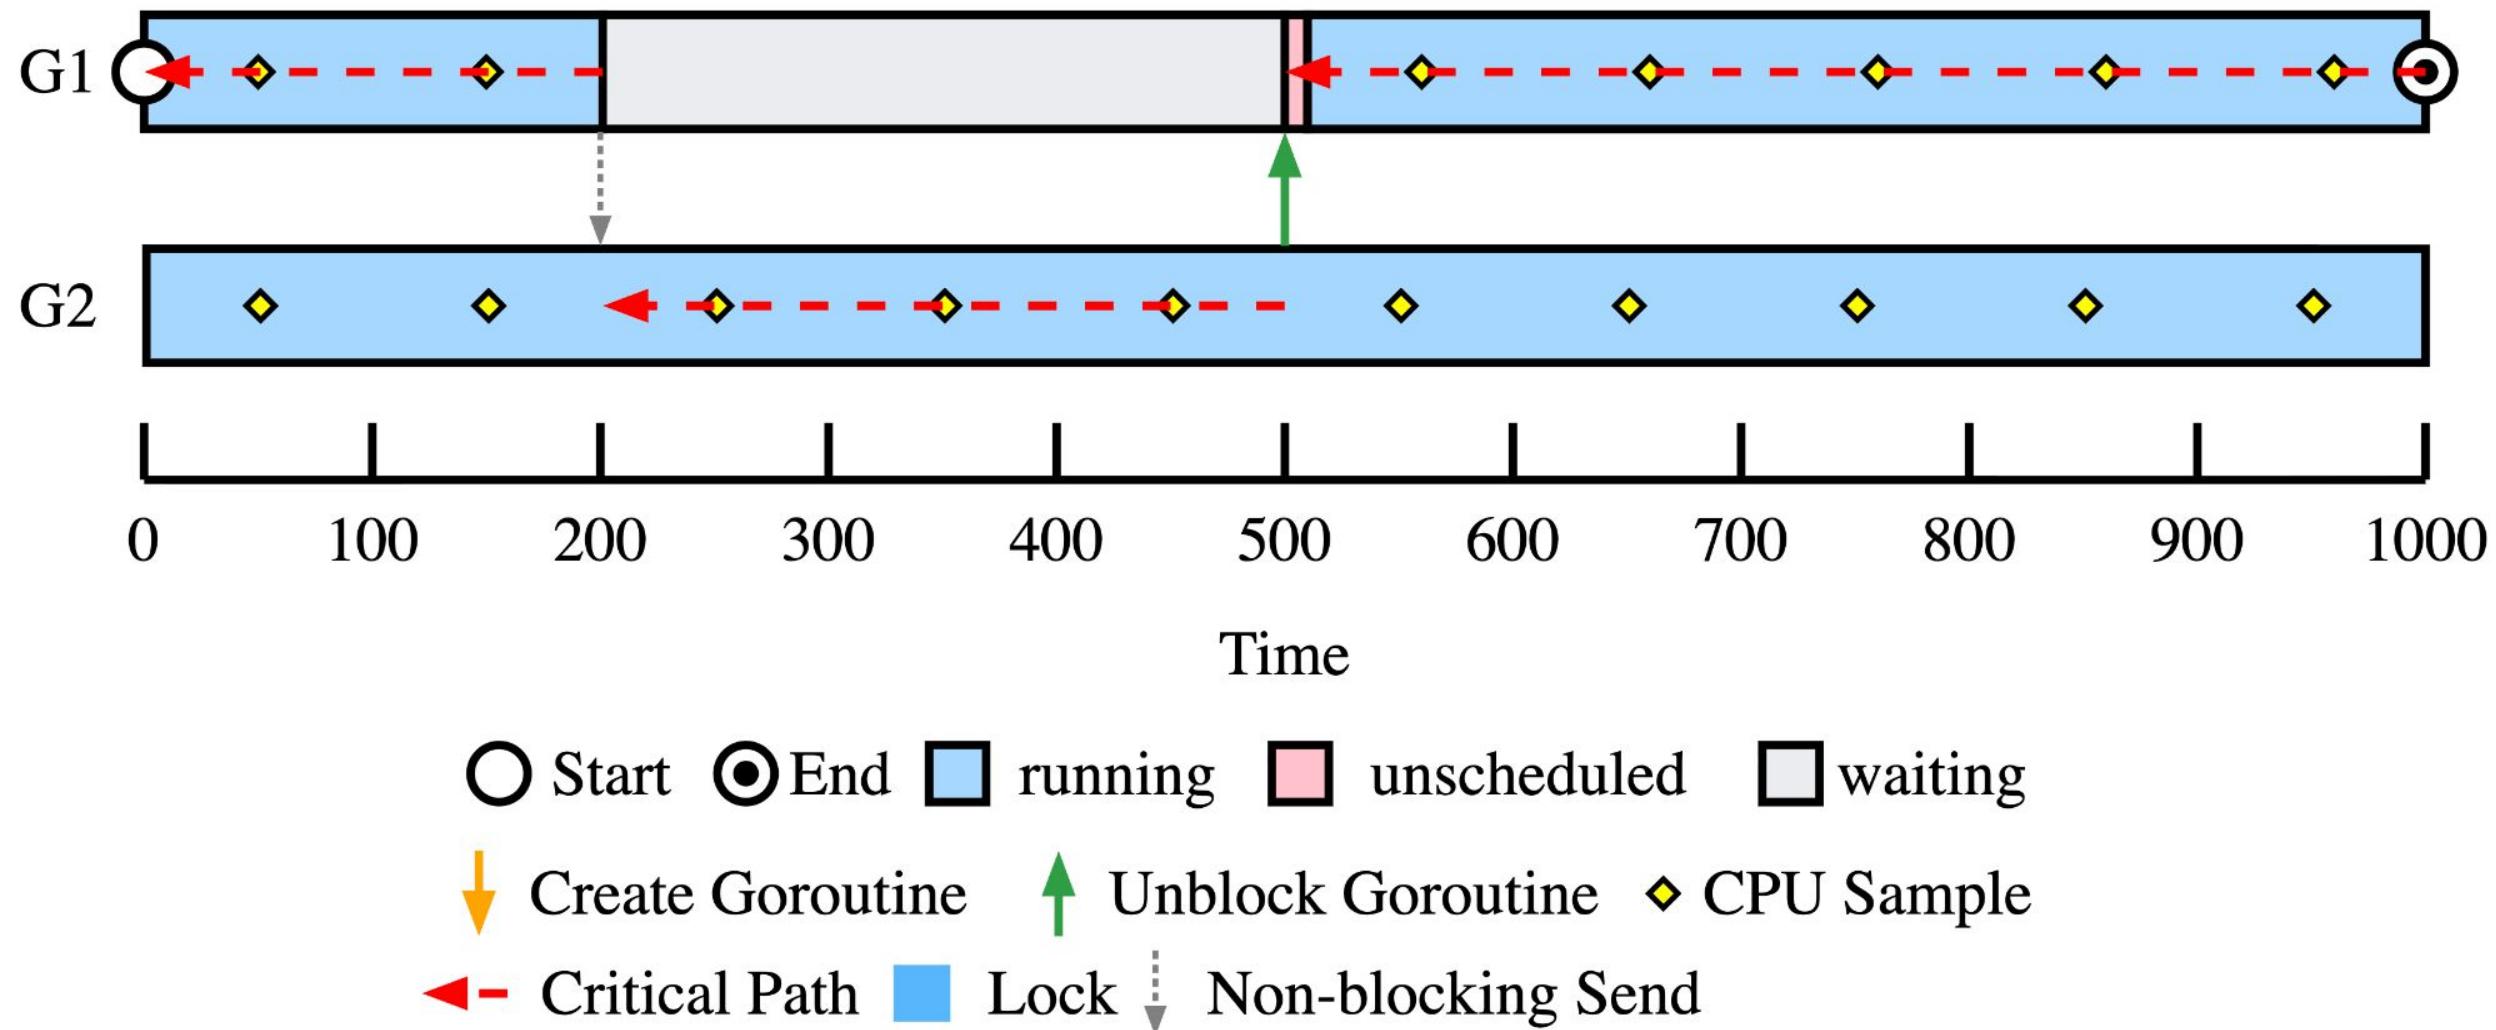


- Start ○ End ■ running ■ unscheduled □ waiting
- ↓ Create Goroutine ↑ Unblock Goroutine ◆ CPU Sample
- ← Critical Path ■ Lock ↓ Non-blocking Send

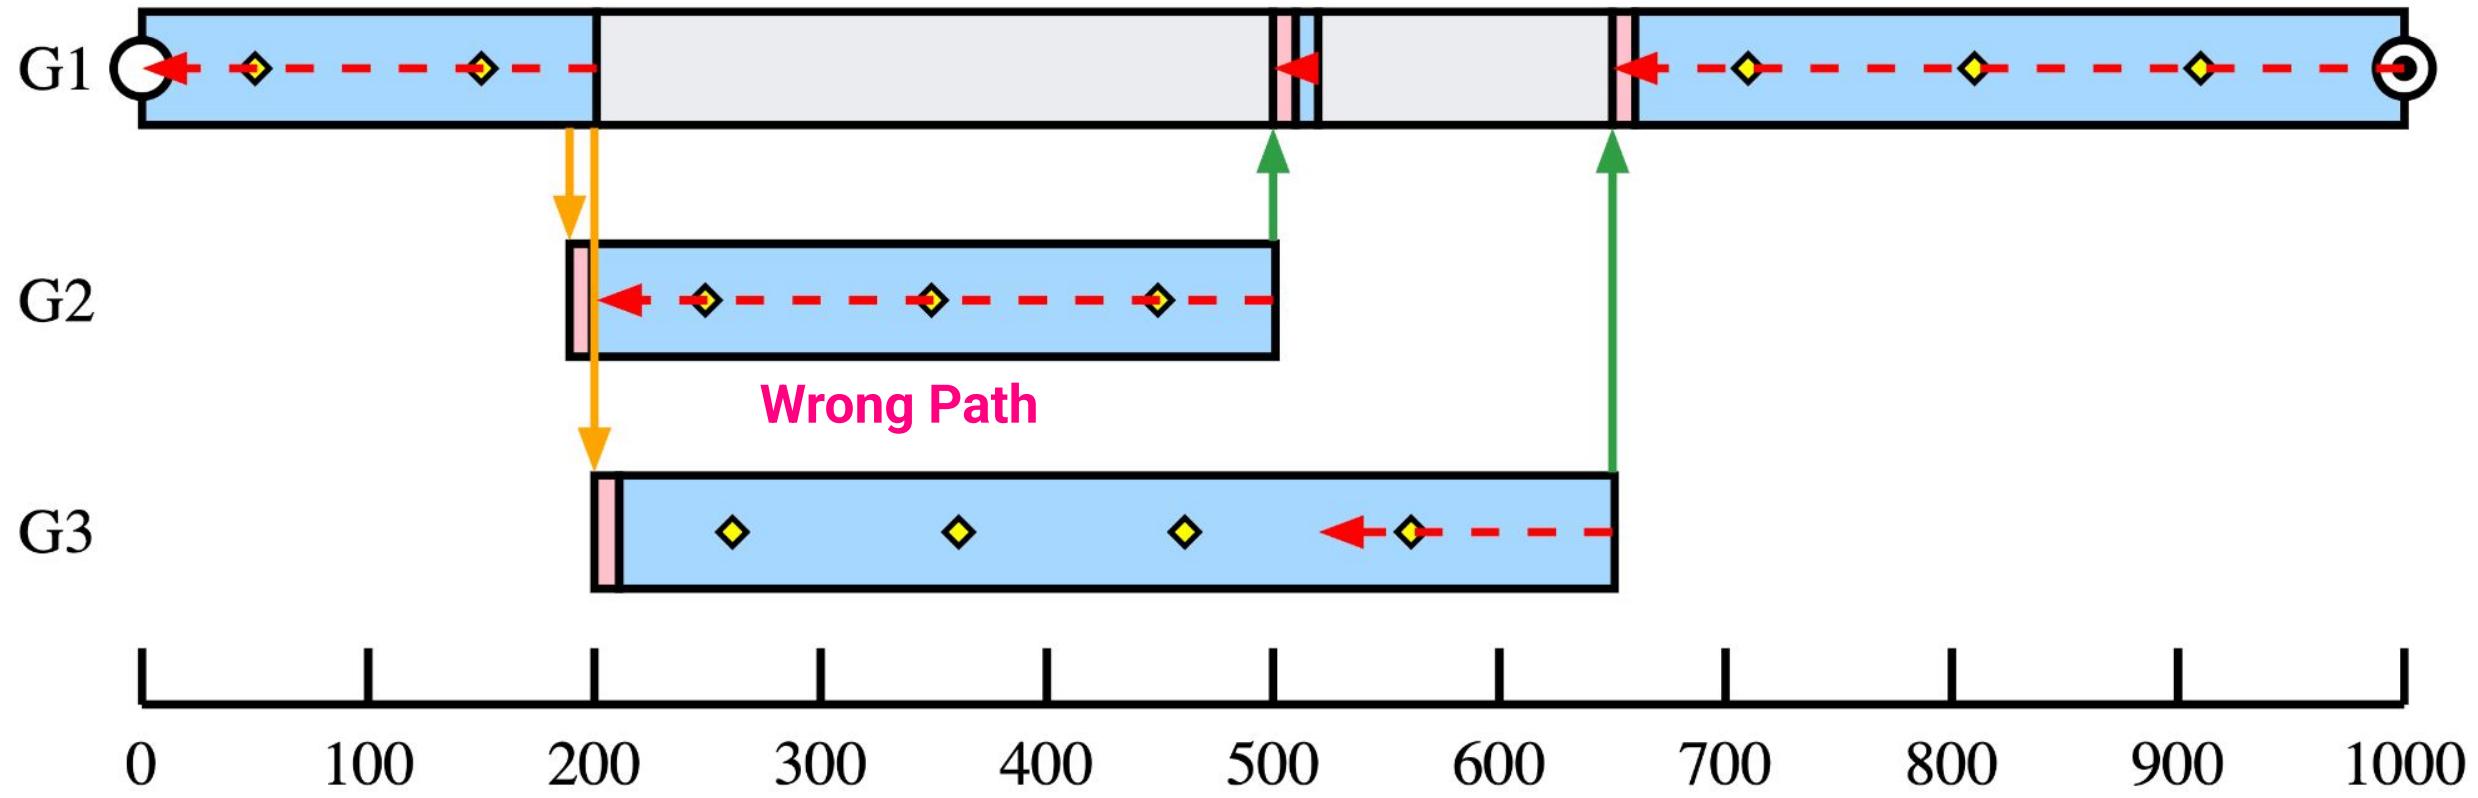
# Worker



# WorkerBusyBuffered



# FanOut

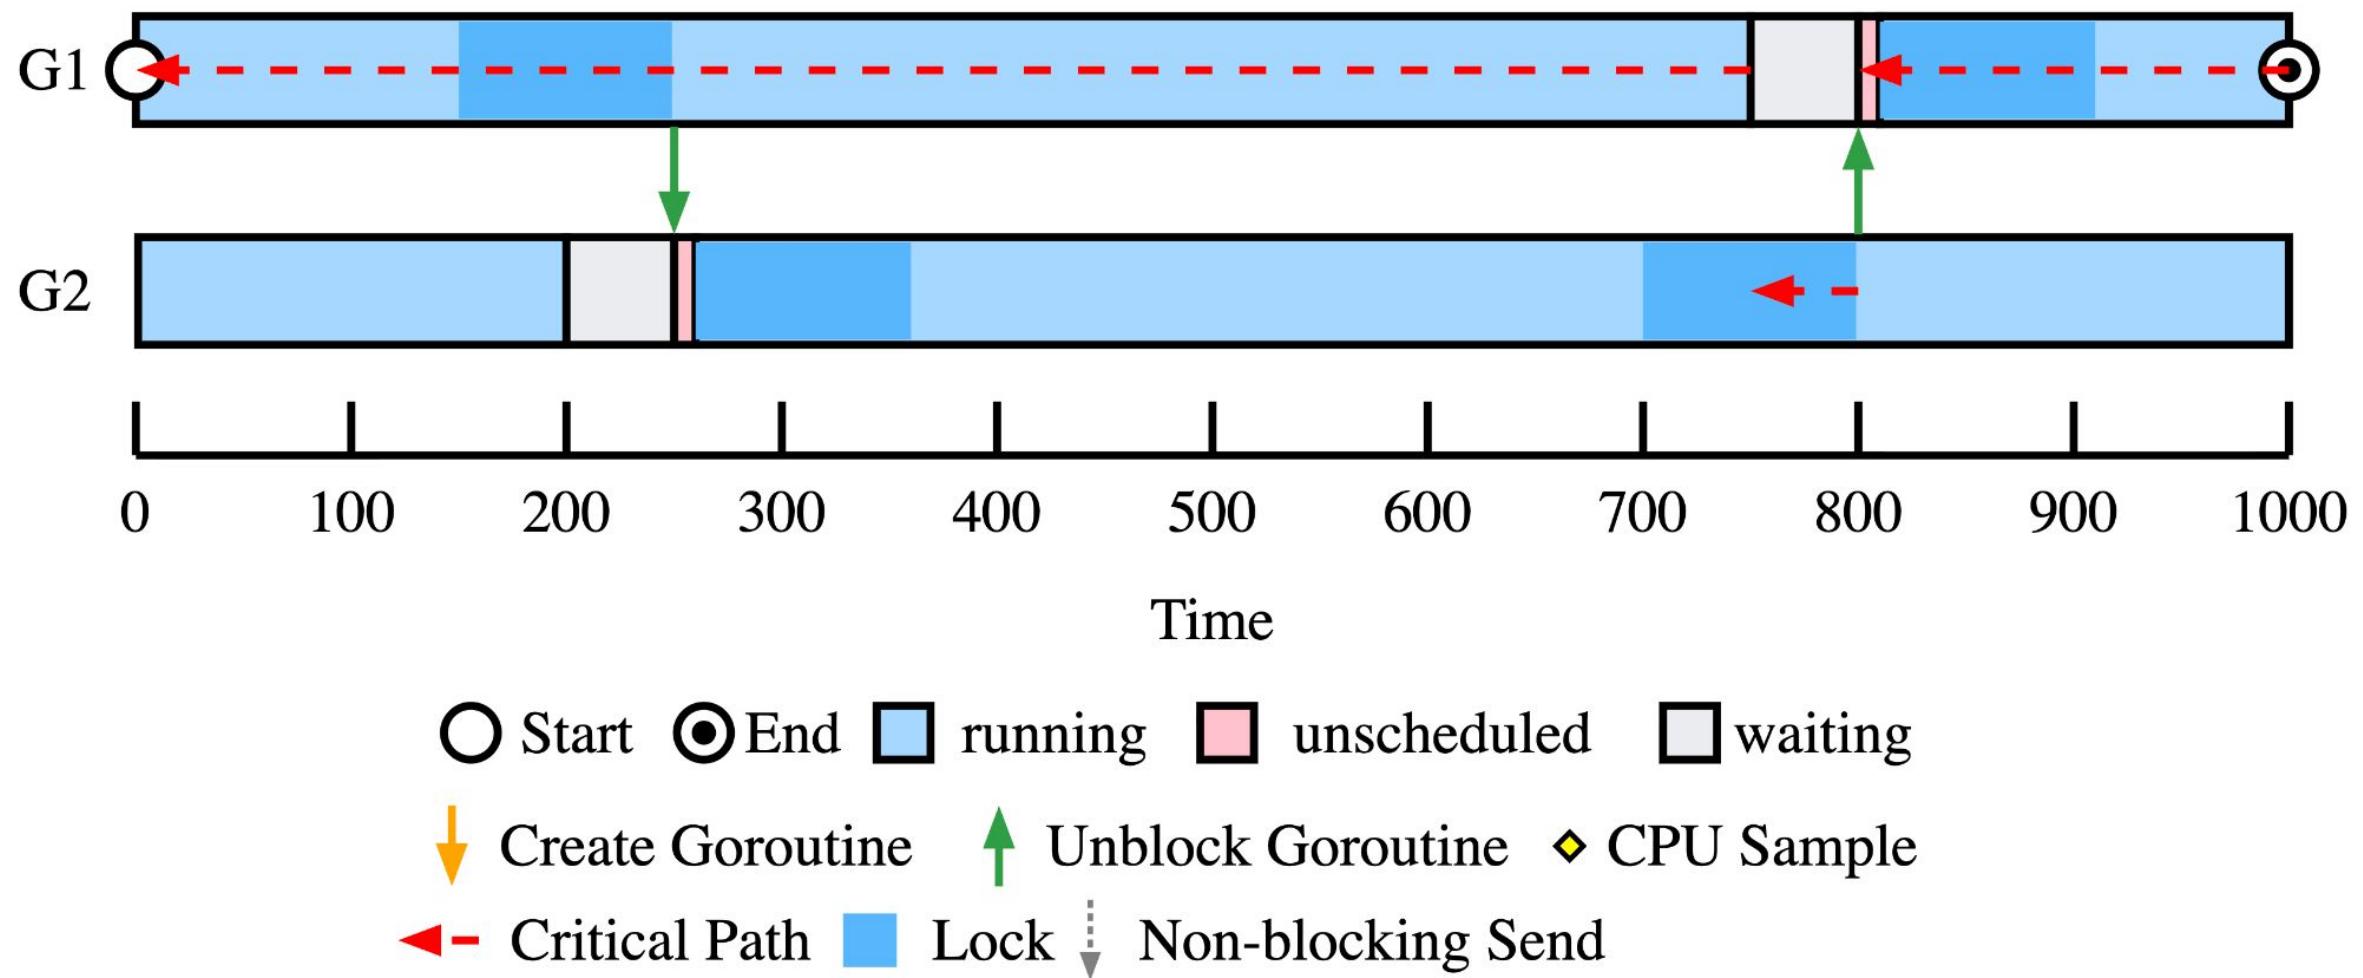


○ Start ○ End ■ running ■ unscheduled □ waiting  
↓ Create Goroutine ↑ Unblock Goroutine ♦ CPU Sample  
← Critical Path ■ Lock ↓ Non-blocking Send

# Critical Path Analysis Algorithm for Go Traces

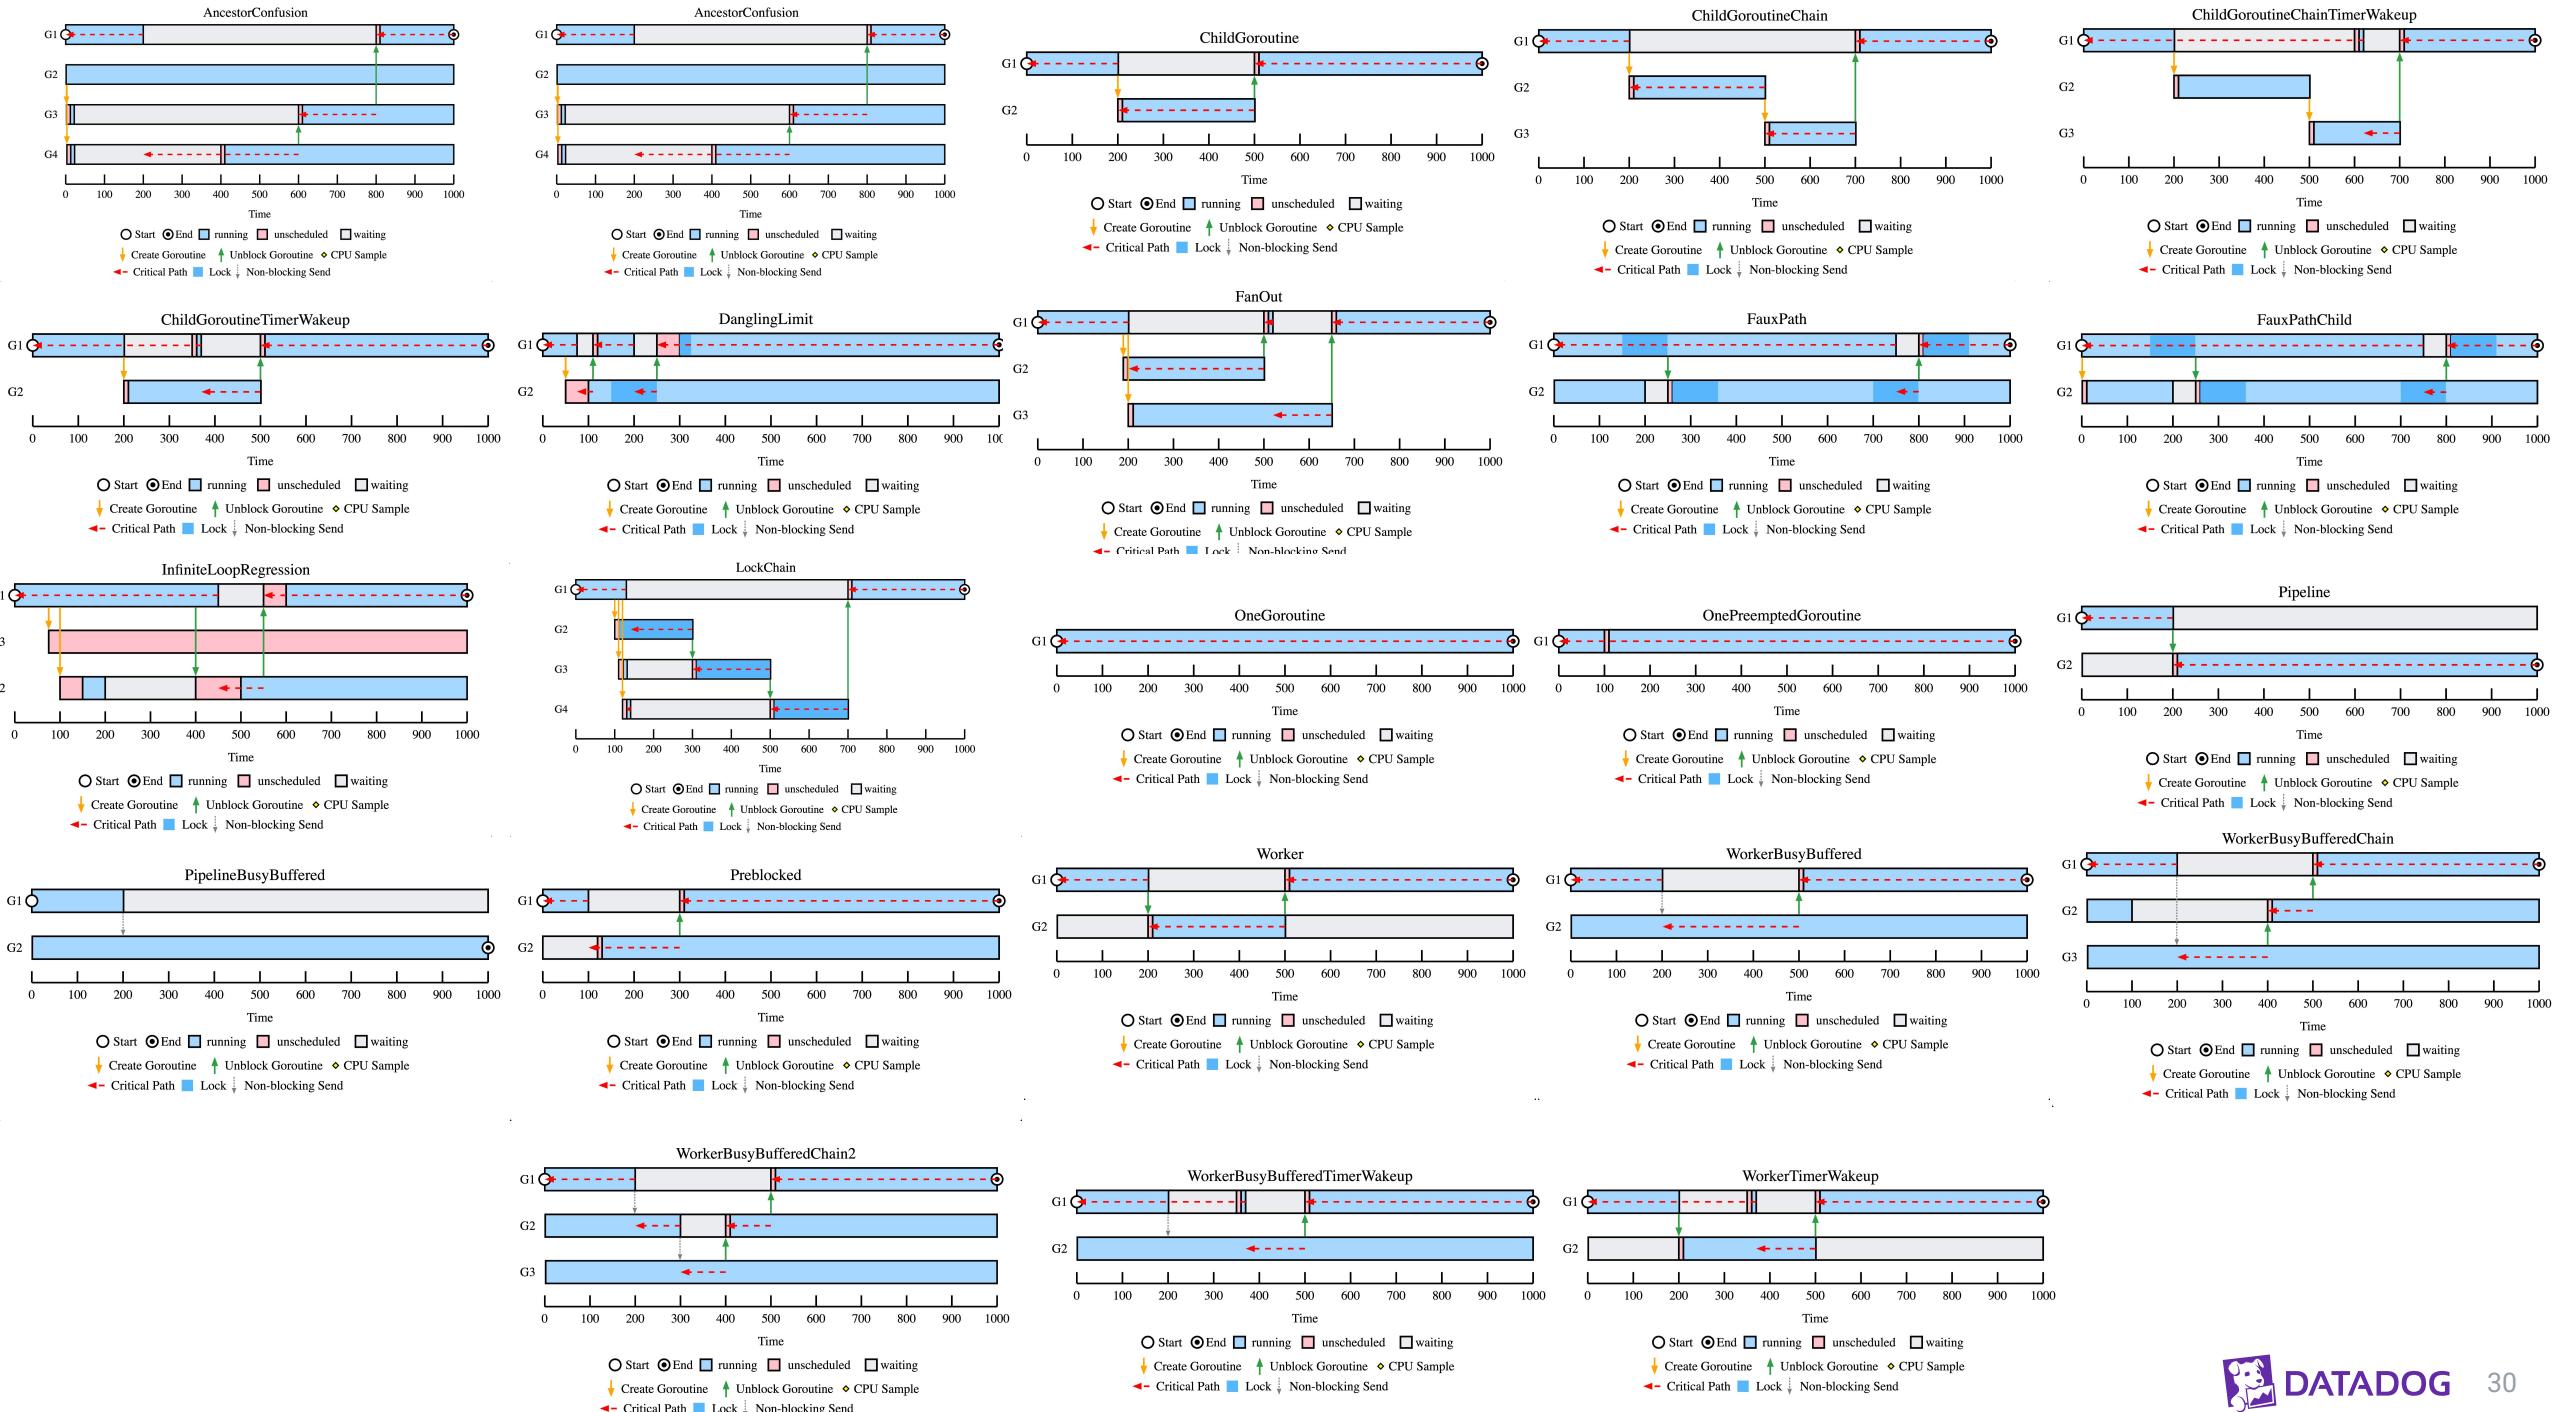
1. Begin the path on the goroutine of the last event.
2. If the event is a wakeup or goroutine creation.
  - a. If the event is a wakeup and the waking goroutine does not have a path to the first event or the stack trace indicates a mutex wakeup:
    - i. Limit the waking goroutine to the start time of the wait preceding the wakeup.
    - ii. If the current goroutine already has a limit, use the max of the two limits.
  - b. Switch to the waking/creating goroutine and continue at step 2.
3. Find the next event mentioning the current goroutine.
4. If the next event is the first event: Success, return the path.
5. If there is a limit and the next event is before it, or there is no next event, return to the goroutine that set the limit.
6. If there is no next event: Failure, return without a path. (*should be impossible*)
7. Continue with the next event at step 2.

## FauxPath

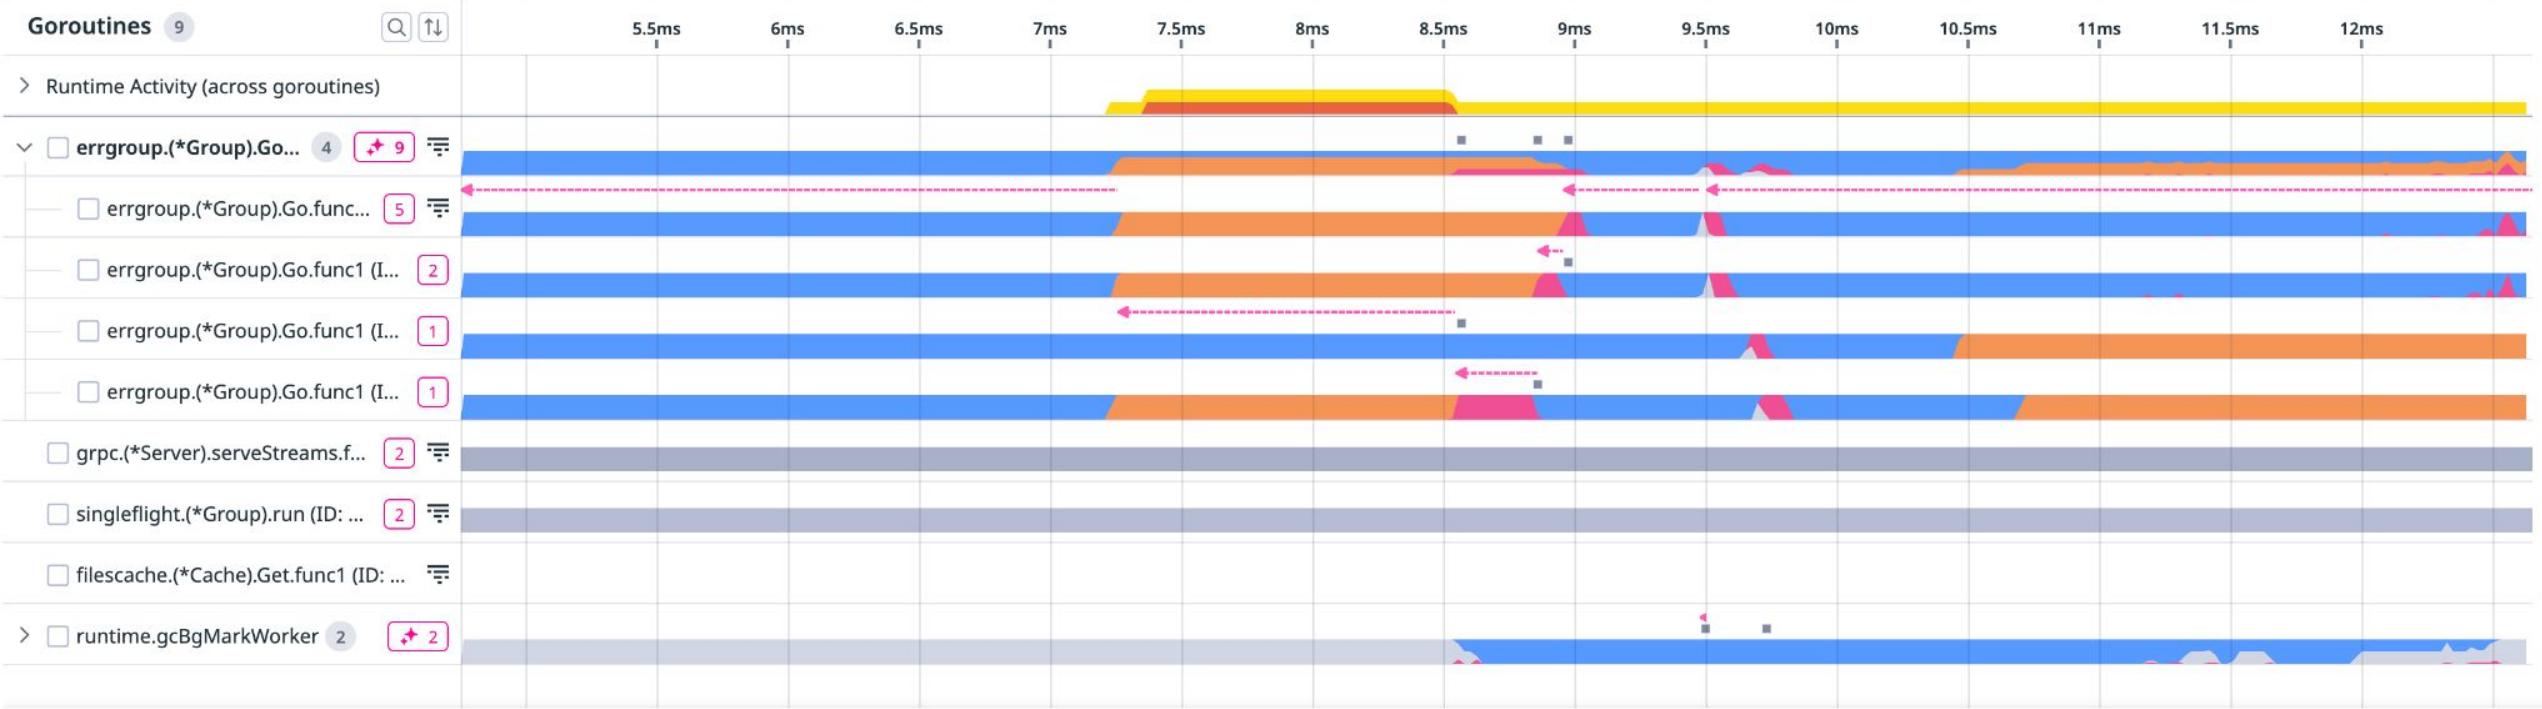


# Critical Path Critique

- Some fundamental issues:
  - Mutual exclusion heuristics based on stack traces are not sound.
  - Lost edges on non-blocking send operations.
  - Lost edges due to OS syscall wakeups (e.g. unix pipes).
- But ... perfection is the enemy of






[Summary](#) [Call Tree](#) [Flame Graph](#)
**Details**

Events 133

Selected Duration 7.89ms

Goroutines All

**Goroutine Activity** 63.14ms**Runtime Activity**

Insight: Critical path This activity was on the critical path for the span to finish

[View Details →](#)

root

errgroup.(\*Group).add.func1

converter.(\*RowConverter).parseRows-fm

converter.(\*RowConverter).parseRows

log.(\*traceLogger).Errorf

log.(\*zapLogger).Error

zap.(\*Logger).Error

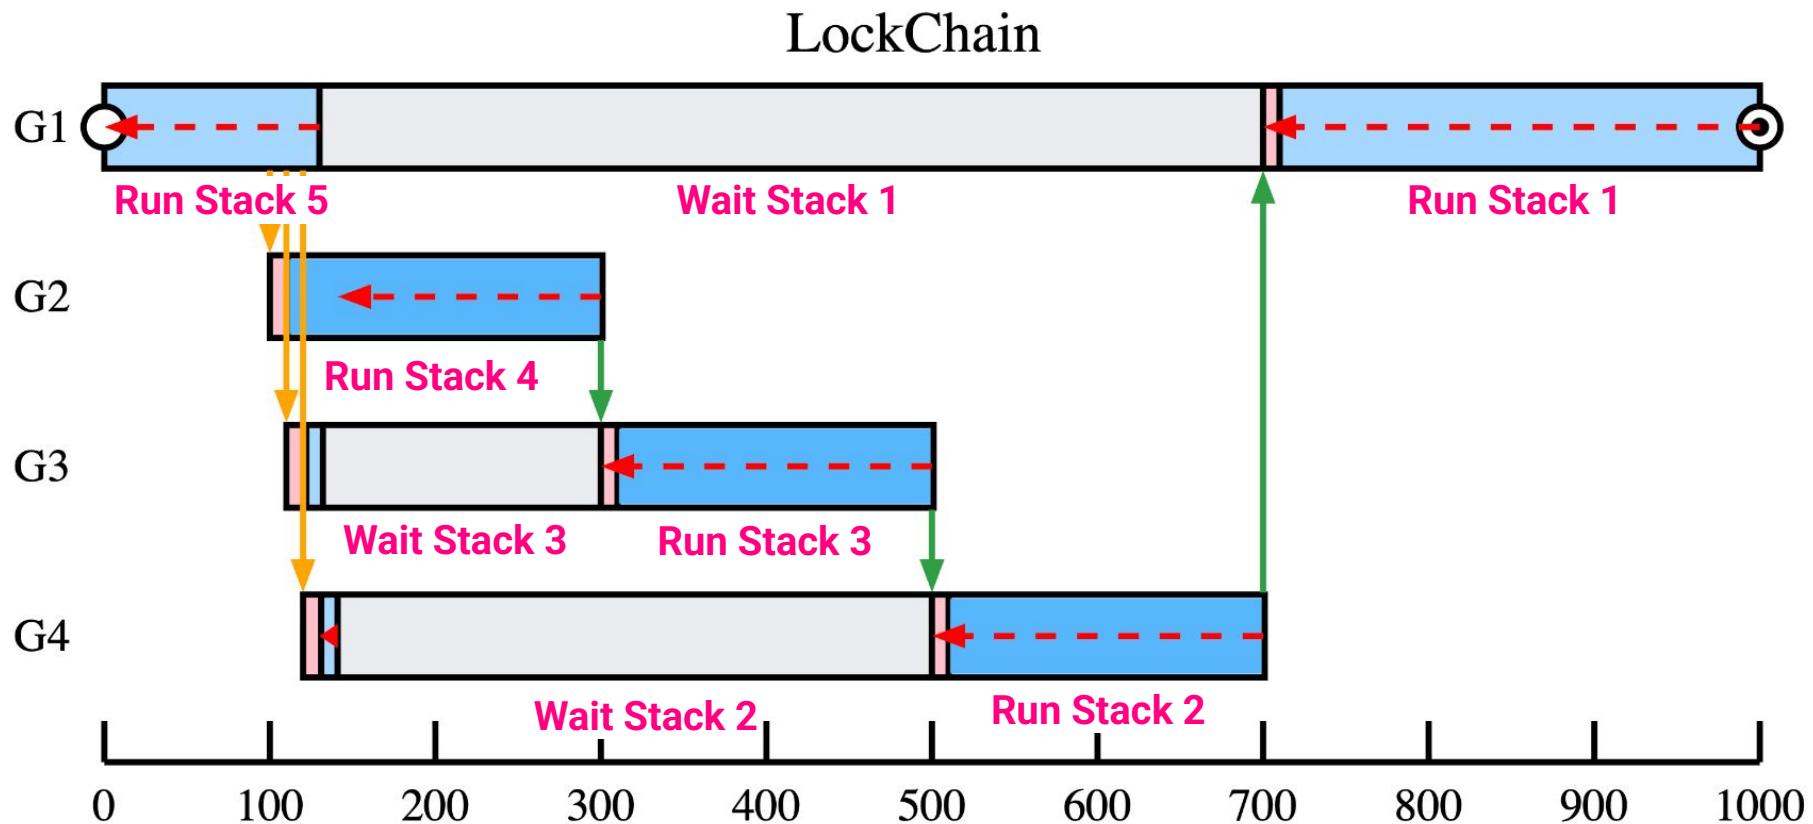
zapcore.(\*CheckedEntry).Write

zapcore.(\*ioCore).Write

zapcore.(\*lockedWriteSyncer).Write

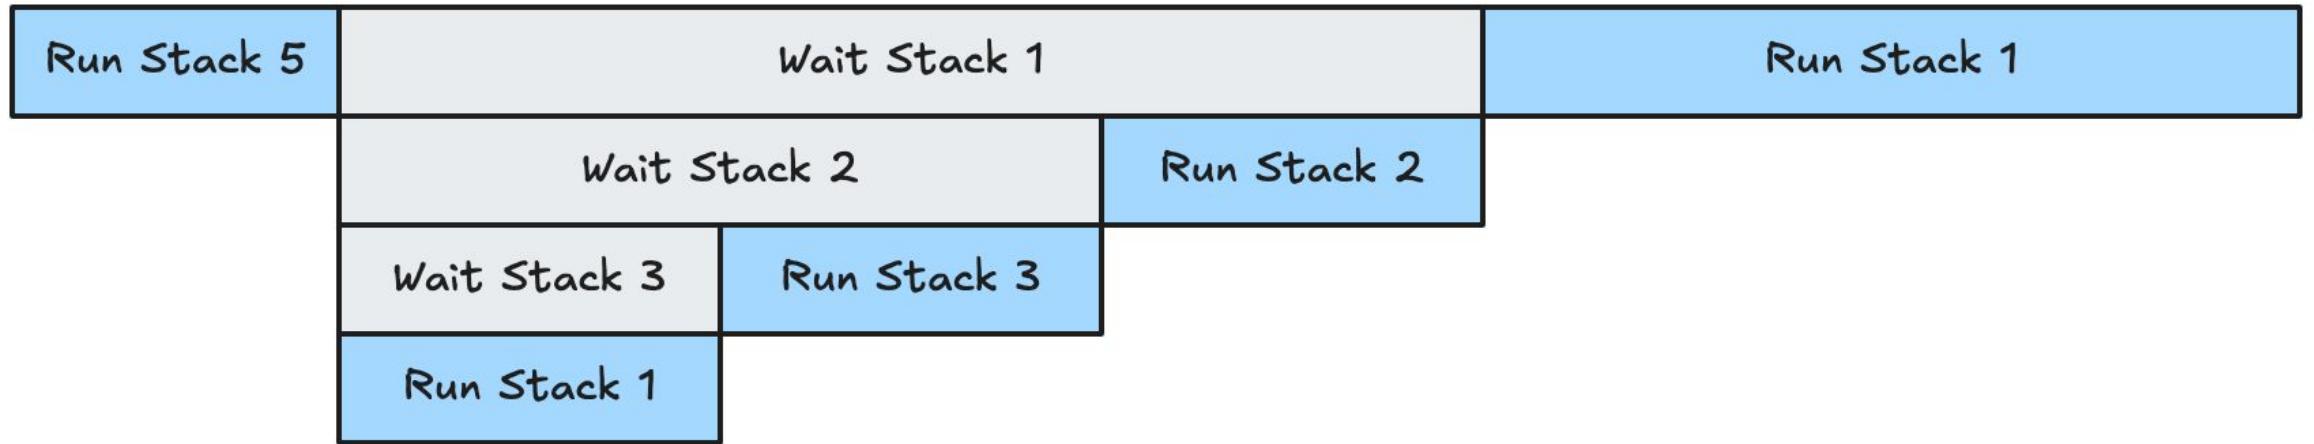
sync.(\*Mutex).Unlock

# Wait Stacks



○ Start ○ End ■ running ■ unscheduled □ waiting  
↓ Create Goroutine ↑ Unblock Goroutine ◆ CPU Sample  
← Critical Path ■ Lock ↓ Non-blocking Send

# Wait Stacks



# Datadog ❤️ Open Source

# Low-overhead tracing (landed in go1.21)



## Low-overhead tracing

Prior to Go 1.21, the run-time overhead of tracing was somewhere between 10–20% CPU for many applications, which limits tracing to situational usage, rather than continuous usage like CPU profiling. It turned out that much of the cost of tracing came down to tracebacks. Many events produced by the runtime have stack traces attached, which are invaluable to actually identifying what goroutines were doing at key moments in their execution.

Thanks to work by Felix Geisendorfer and Nick Ripley on optimizing the efficiency of tracebacks, the run-time CPU overhead of execution traces has been cut dramatically, down to 1–2% for many applications. You can read more about the work done here in [Felix's great blog post](#) on the topic.

# CL Pending: Create traces for testing

x/exp/trace,internal/trace: support event constructor for testing #74826

Edit

New issue



Open



felixge opened last month · edited by felixge

Edits

Contributor

...

As discussed in [#62627](#), this issue proposed to extend the [API of x/exp/trace](#) to support programmatically creating trace events for testing and analysis purposes.

*Note: Given that the code for this lives in `internal/trace` right now and is only exposed via `x/exp/trace`, I think this issue doesn't need to go through the official proposal process.*

## Overview

The proposal is to add an API using a `MakeEvent` constructor that takes a `EventConfig` struct. Below is an example for creating a goroutine state transition event:

```
details := MakeGoStateTransition(5, GoRunning, GoWaiting)
details.Stack = MakeStack([]StackFrame{
    → {PC: 1, Func: "time.Sleep", File: "runtime.go", Line: 10},
    → {PC: 2, Func: "main", File: "main.go", Line: 20},
})
details.Reason = "sleep"
ev, err := MakeEvent(EventConfig[StateTransition]{
    ... Time: ..... timestamp,
    ... Goroutine: 1,
    ... Proc: ..... 2,
    ... Thread: ..... 3,
    ... Details: ..... details,
})
```

### Assignees

No one assigned

### Labels

FeatureRequest LibraryProposal

NeedsDecision

### Type

No type

### Projects

No projects

### Milestone

Unreleased

Due by December 31, 2099, 82% complete

### Relationships

None yet

### Development

Code with agent mode

No branches or pull requests

A partially completed prototype of this API is available in [CL 691815](#).

# Thank You

Follow @felixge.de (BlueSky) or, felixge (X) for updates on Critical Path Analysis